

# Web: HTTP (*HyperText Transfer Protocol*) and other related protocols

## Telematics Applications

Bachelor in Telecommunication Technologies Engineering

Slides generated by Florina Almenarez, Daniel Díaz, Andrés Marín, Patricia Callejo  
Reproduction is forbidden without the permission of the authors except for University Carlos III students. Images from books (if any) belongs to their corresponding authors

apitel - GitLab

<https://gitlab.gast.uc3m.es/apitel>

apitel

Group ID: 110

Subgroups and projects

Project	Created
2016-2017-assignment	5 years ago
2017-2018-assignment	4 years ago
2018-2019-assignment	3 years ago
2021-2022-assignment	10 months ago
dns	11 months ago
email	
http	
intro-sockets	
lab	
optional-dns-recon	
optional-java-socket-server-example	
optional-tcp-advanced	

# Index

## General perspective

1. Introduction
2. First versions of HTTP
3. HTTP 1.1
4. HTTP/2
5. Other protocols

## Bibliography

- **Basic:**

- Forouzan, Behrouz A.. TCP/IP Protocol Suite. McGraw-Hill Higher Education. 2006 (Capítulo 22).

- **Complementary:**

- “Internetworking with TCP/IP Volume I. Principles, Protocols and Architecture”, 5a Ed. D.E. Comer and D.L. Stevens, Prentice-Hall Int., 2006 (Capítulo 28).
- “Learning HTTP/2: A Practical Guide for Beginners”, Stephen Ludin & Javier Garza. O’reilly, June 2017.

- **RFCs:**

- RFC 2616: Hypertext Transfer Protocol -- HTTP/1.1, 1999
- RFC 9113: HTTP/2, 2022
  - RFC 7540: Hypertext Transfer Protocol Version 2 (HTTP/2), 2015
- RFC 9114: HTTP/3, 2022

## Objectives

- Know advanced aspects of the most **popular Internet services**: SMTP, **HTTP**, etc.
- Understand the basic concepts and components related to the WWW and its architecture.
- Learn about HTTP and the differences between the different versions.
- Define the fields of an HTTP **request** and **response**
- Define non-persistent and persistent **connections** in HTTP
- Understand Web **caching** and **proxying**
- Learn about **cookies** and the usage in HTTP
- Learn how to make HTTP requests using command line tools, e.g., telnet, curl...

## Content >> Introduction

### 1. Introduction

#### A. Web

#### B. Definitions

#### C. General operation

### Index

#### 1. Introduction

#### 2. First versions of HTTP

#### 3. HTTP 1.1

#### 4. HTTP/2

#### 5. Other protocols

# HTTP – Hypertext Transfer Protocol

- # HTTP – Hypertext Transfer Protocol

## Introduction &gt;&gt; Web Content

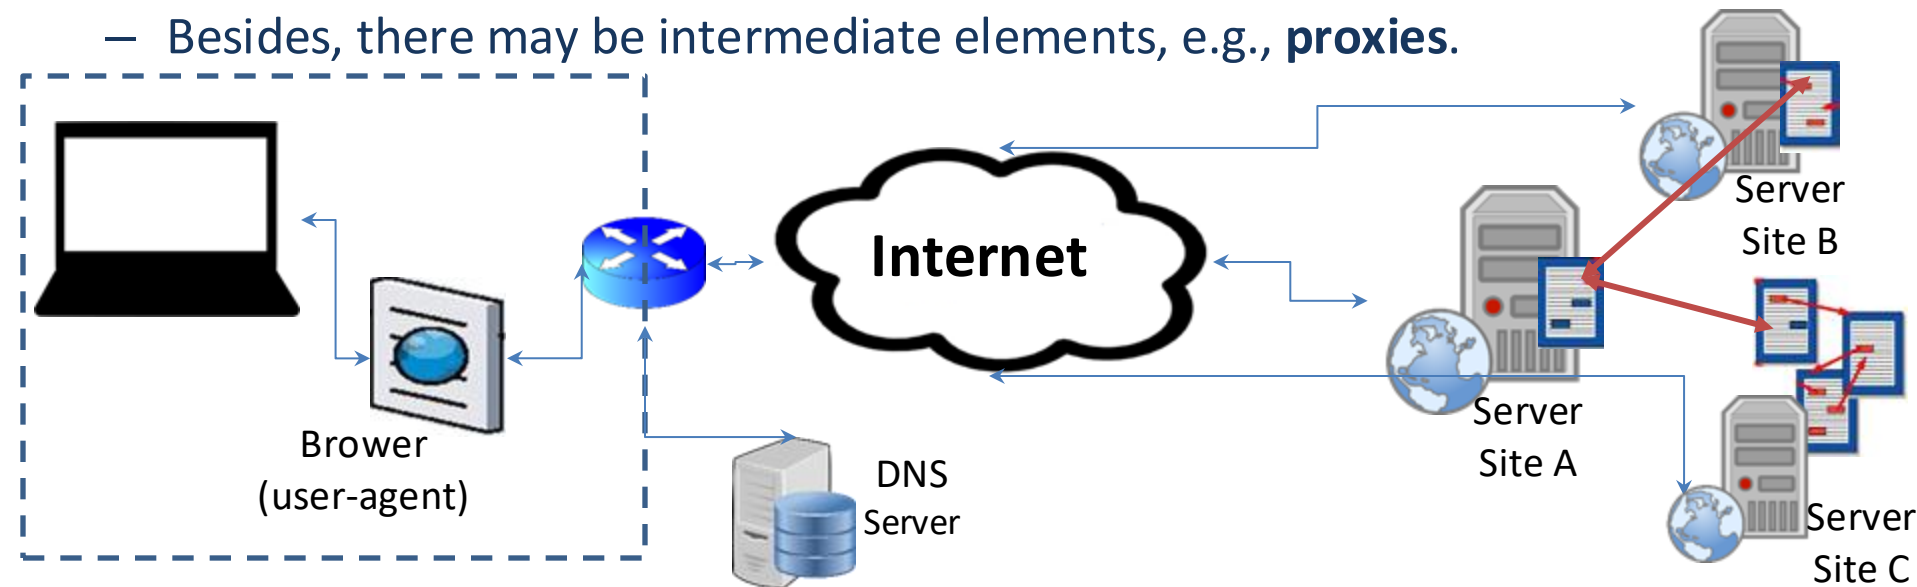
- The content can be grouped into two main categories:
  - **Static**
    - its content does not change, e.g., documents HTML, XML, XSL, XHTML, CSS, etc.
  - **Dynamic**
    - Execute some code (script) on the **server** side, e.g., CGI, Java, Python, PHP, ASP, etc.
    - Execute some code (script) on the **client** side, e.g., Javascript, AJAX, etc. Also known as **active**.
- The websites can have embedded resources (e.g., small base64 images, scripts...).





## Content &gt;&gt; WWW Architecture

- The WWW is a distributed service **client-server**.
  - **client** uses a browser (or other type of program, e.g spiders) to access a **service** provided by a **server**.
  - the service is distributed over different locations called "**sites**".
  - each site hosts one or more hypertext documents (resources), called "**web pages**".
  - each web page may contain one or more links to other web pages
  - Besides, there may be intermediate elements, e.g., **proxies**.

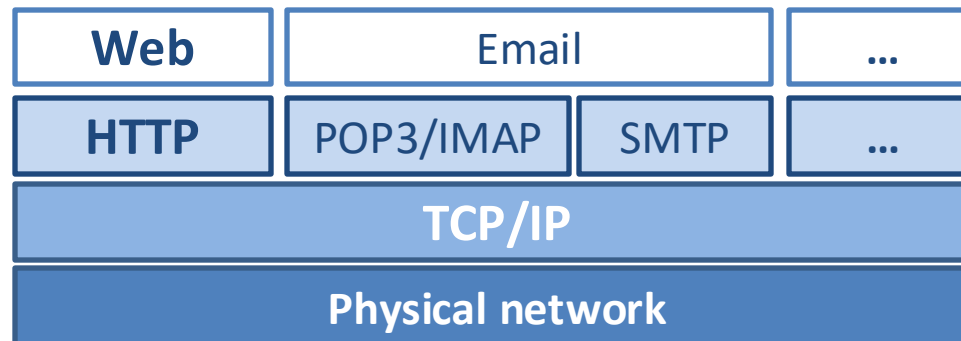


## Introduction &gt;&gt; Definitions

- *Hypertext Transfer Protocol (HTTP)*  $\Rightarrow$  TCP/IP application layer protocol for distributed, collaborative hypermedia information systems [IETF, RFC2616, 1999]
  - Invented by Tim Berners-Lee, W3C, 1989-1991
- Allows a client's device (computer, smartphone, household appliances, etc.) request a server to send a resource....
  - **hypertext/hypermedia** transfer protocol
  - nowadays, it can be used for other tasks, e.g. in distributed object management systems
  - allows systems to be built independently of the data being transferred.
- The protocol allows communicating servers, proxies, gateways and clients through the **request/response** scheme.
  - it is based on plaintext messages
  - readable, easy to debug

## Introduction &gt;&gt; Definitions


- **Connection-oriented** and **stateless** protocol
  - **connection-oriented**, it uses the **TCP** transport protocol for its operation (end-to-end between a client and a Web server)\*.
  - **port 80**, default for non-secure connections (`http://`)
  - **port 443**, default for secure connections (`https://`)



- **No maintains the state** after each request-response pair, as each connection is independent of the previous one.
  - there is no "session" concept.


\* Other transport protocols are currently supported

## Introduction >> General Operation




- The user types in the address bar of the browser the resource (URL) he/she wants to access:
  - <http://www.uc3m.es/>
- The browser breaks the URL into 3 parts:
  - **Scheme:** protocol ("`http`")
  - **Authority:** name of the server/domain ("`www.uc3m.es`")
  - **Path to the resource:** ("`/`")
- The browser communicates with the DNS to translate the name "`www.uc3m.es`" into an IP address, which is used to connect to the Web server machine (176.58.10.138).

## Introduction &gt;&gt; General Operation



- The Web server, listening on port 80, accepts the connection.
- Client sends the **request**, e.g.: `GET /`
- The HTTP server receives the request message, creates a response message including the HTML text of the requested page (e.g. root resource) and **send the response**.
- The HTTP server **closes\*** the TCP **connection** (depends on the version).
- The client receives the response message, and renders the web page in the browser.

## Introduction >> General Operation



How were or are the request and response messages?

How would the full page be retrieved if it contains 20 images that are part of the page but not "embedded" in it?



*Depends on the protocol version*

**HTTP 0.9**

**HTTP 1.0**

**HTTP 1.1**

**HTTP/2**

**HTTP/3**

## Content >> First versions

### 2. First HTTP versions

A. HTTP 0.9

B. HTTP 1.0

C. Exercise

### Index

1. Introduction

**2. First HTTP versions**

3. HTTP 1.1

4. HTTP/2

5. Other protocols

## First versions >> HTTP 0.9

- Named after the original version, defined in 1991, W3C
- The **request** consists of a single line of ASCII characters, starting with the **GET** method, followed by the resource to be requested and ending with <CRLF>, e.g.,

```
GET /mywebpage.html
```

- The **response** consists of the requested resource ("*byte stream*" of ASCII characters) and ends with the closing of the connection by the server, e.g.,

```
<html>
```

```
A very simple and obsolete web page
```

```
</html>
```

- One connection is used for each resource

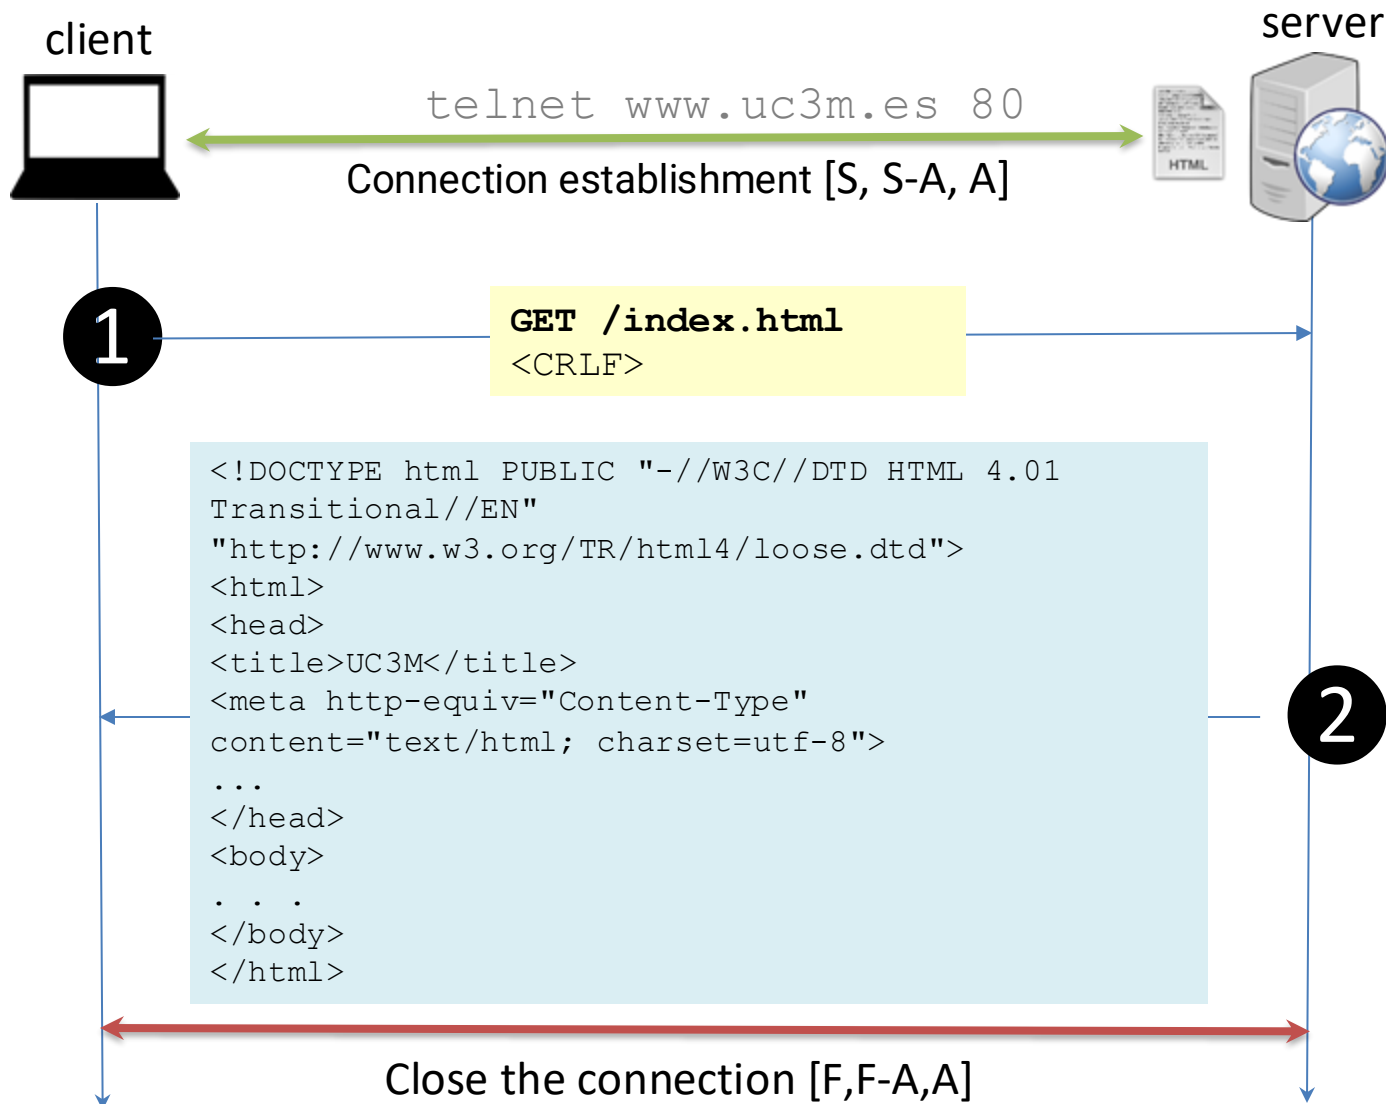
No HTTP headers

Only HTML files could be transmitted

No response status information or error codes



## First versions >> Example HTTP 0.9



## First versions >> HTTP 1.0

- HTTP/1.0, RFC 1945, May 1996
- The protocol version is sent in each request
- Different methods are defined for the request: GET, HEAD, POST
- The concept of headers is added, both for requests and responses.
- Allows the use of different types of resources
  - messages are exchanged in a format similar to those used by RFC 822 and **MIME** (*Multipurpose Internet Mail Extensions*)
- A status code is sent at the beginning of the response.

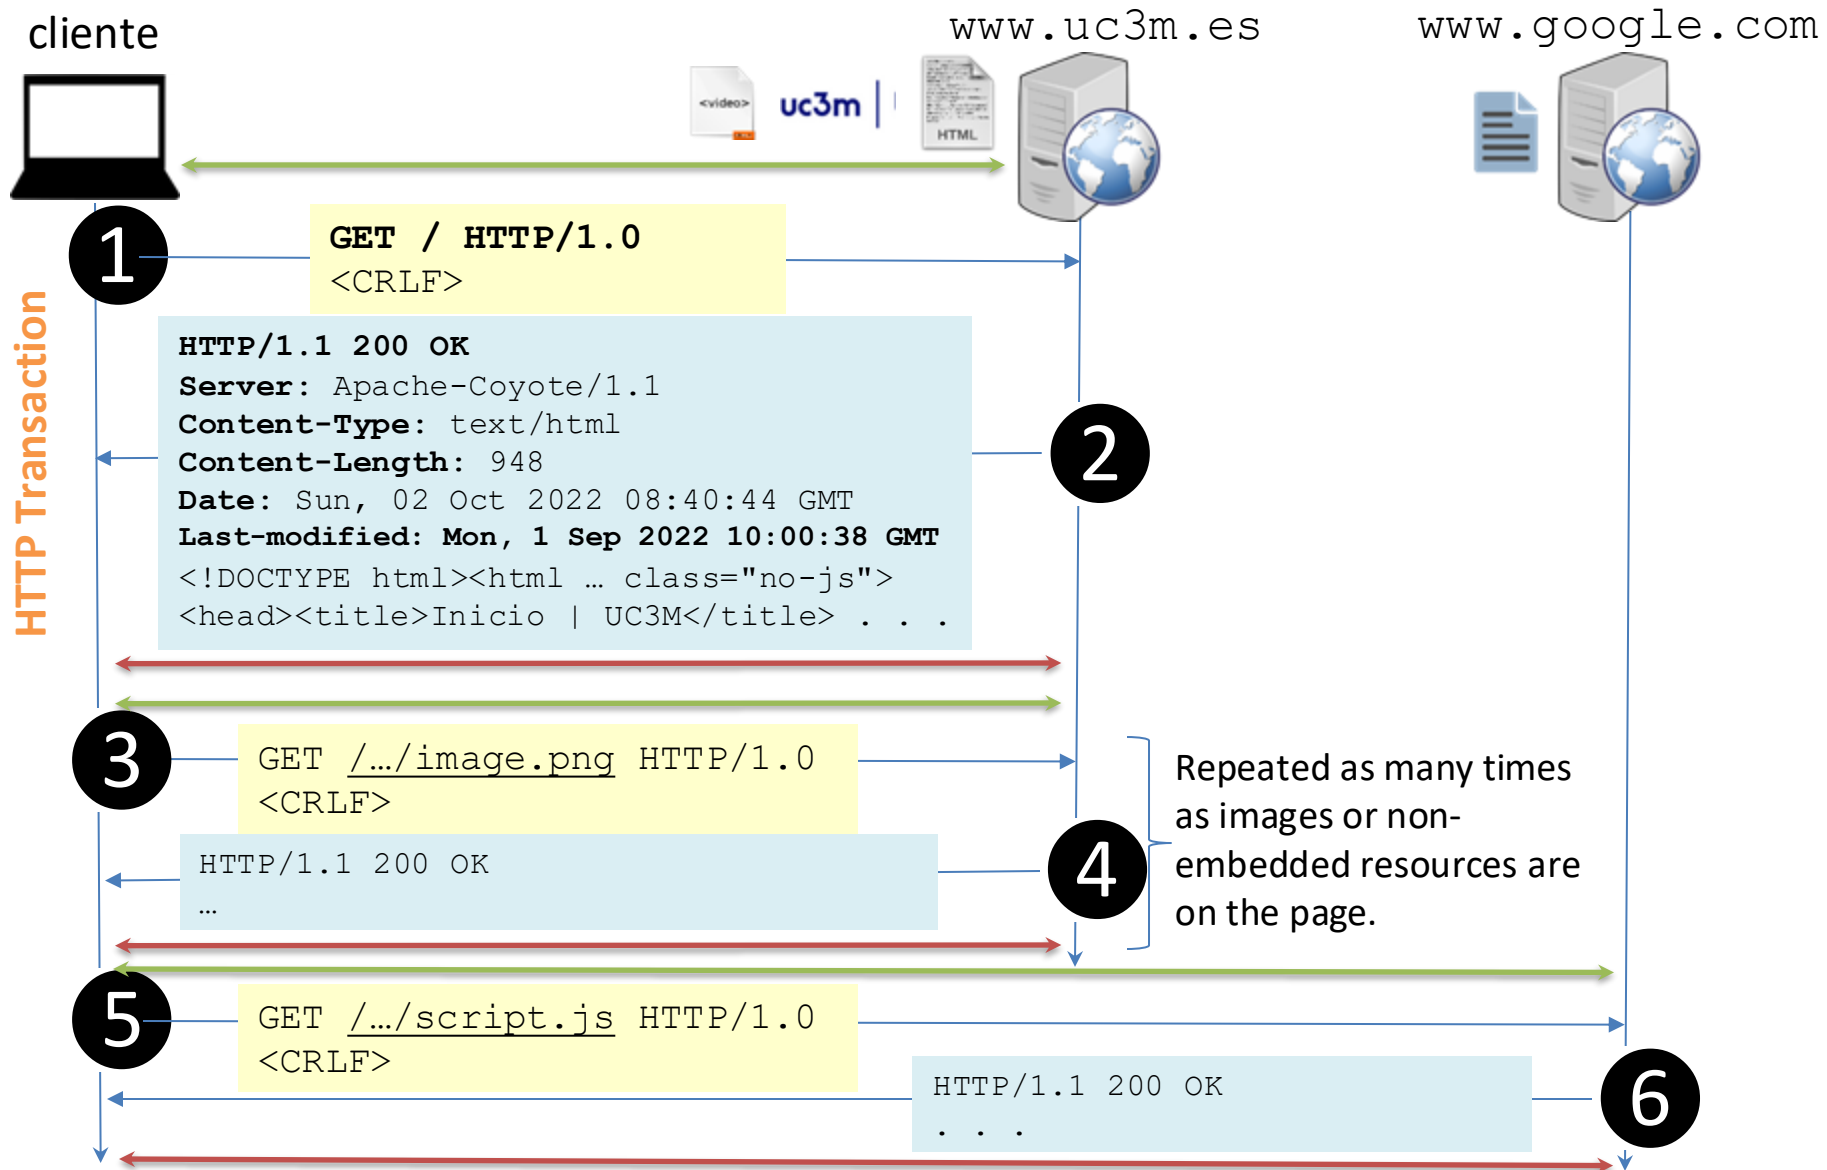
### Full request

```
Method Request-URI HTTP-Version CRLF
* ( General-Header
  | Request-Header
  | Entity-Header )
CRLF (an empty line)
[ Entity-Body ]
```

### Full Answer

```
HTTP-Version Status-Code Reason-Phrase CRLF
* ( General-Header
  | Response-Header
  | Entity-Header )
CRLF (an empty line)
[ Entity-Body ]
```

## First versions >> Example HTTP/1.0



## First versions >> Exercise

- What is the difference between request-response messages in HTTP/1.0 compared to HTTP/0.9?



- What type of documents can you retrieve with HTTP/0.9?



- What is an HTTP transaction?



## Content >> HTTP 1.1

### 3. HTTP 1.1

- A. General information
- B. Connection types
- C. Proxy and cache
- D. Request-response messages
- E. Persistent connections
- F. Head of line blocking
- G. Status management
- H. Exercise


### Index

- 1. Introduction
- 2. First HTTP versions
- 3. HTTP 1.1**
- 4. HTTP/2
- 5. Other protocols


## HTTP 1.1

- HTTP/1.1, RFC 2616, June 1999
- Maintains features of HTTP/1.0, e.g., use of MIME, specify protocol version, response status code.
- Extends other features such as **methods** (optional) and request and response **headers** (from 16 in HTTP/1.0 to 46 in HTTP/1.1).
- Takes into account the need for **persistent connections**, virtual hosts and the effects of hierarchical proxies and **cache** usage.
  - A **proxy** is an intermediary program that acts as a server to the client or as a client to the server for the purpose of making requests on behalf of other clients.
  - **Cache** is the local store of response messages and the subsystem that controls the storage, retrieval and deletion of messages.
    - reduce response times and network bandwidth consumption on future requests

## HTTP/1.1 >> Client-Server connection types



### Separate connections / HTTP request (absolute URI, p.ej. GET <http://www.w3.org/index.html> HTTP/1.1)



## HTTP/1.1 >> *Proxies*

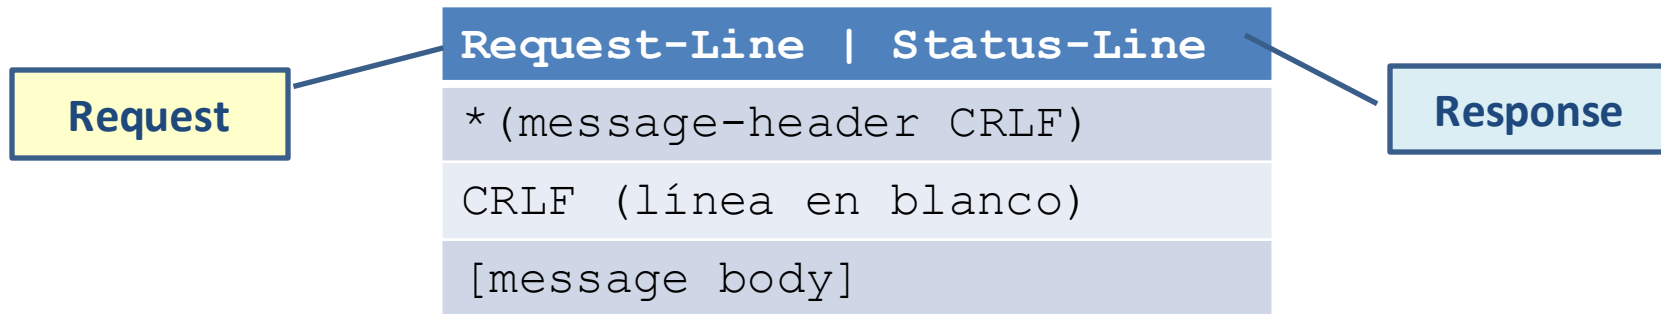
- Uses:
  - **Security**
    - Client > proxy > firewall > Internet
    - proxy: machine authorized to make HTTP connections
  - **Transform** requests and responses
    - translates HTTP versions
    - compress responses to suit low-capacity clients
  - **Establish a **cache**** of requests and responses
    - prioritize proxy traffic
    - save bandwidth (BW)
    - reduce user-perceived latency



## HTTP/1.1 >> Proxy-cache

- Explicit use:
  - users configure their browsers to redirect their access through a proxy-cache, rather than to the end server hosting the resources they are seeking
  - drawbacks:
    - browser configuration required
    - if the proxy is not available, the browser must be reconfigured.
- Transparent:
  - perimeter systems intercept and redirect requests  $\Rightarrow$  TCP SYN lookups (without ACK) to ports 80
  - located at the edge of the ISP (*“Internet Service Provider”*)
  - users are unaware of their existence
- It is essential to have the resources updated in cache
  - `expiry` or validation mechanisms
  - time to modify, time in cache, etc.
  - cache parameters in HTML (HTTP/1.0 y 1.1)

# HTTP/1.1 >> Request-Response Messages



**Request-Line** = **Method** Request-URI HTTP-version CRLF

## Methods

OPTIONS	Check the available options
GET	<b>Request a recourse (Request-URI) to the server</b>
HEAD	Requests information about a resource, without body (only headers = metadata)
POST	<b>Sends data from the client to the server</b>
PUT	Creates/updates/replaces a document under the provided URI
DELETE	Deletes a web page (Request-URI)
TRACE	Performs a message loop test along the path to the requested resource
CONNECT	Reserved for proxies

# HTTP/1.1 >> Response message

**Status-Line** = HTTP-Version **Status-Code Reason-Phrase** CRLF

Code	Reason-phrase	Description
<b>INFORMATIONAL - 1XX</b>		
100	<i>Continue</i>	Partial response to an incomplete request that should be continued
<b>101</b>	<b>Switching Protocols</b>	<b>The server is willing to change the protocol to the one requested by a client (Upgrade header).</b>
<b>SUCCESS – 2XX</b>		
<b>200</b>	<b>OK</b>	<b>The request has been successful</b>
201	<i>Created</i>	New resource has been created as a result of successful request
202	<i>Accepted</i>	Request received
203	<i>Non-Authoritative Information</i>	Content was not obtained from the originally requested source
204	<i>No Content</i>	The response has no content
205	<i>Reset Content</i>	The UA must initiate the page from which the request was made.
206	<i>Partial Content</i>	Requested content partially served
<b>REDIRECTION – 3XX</b>		
300	<i>Multiple Choices</i>	The request has more than one possible response
<b>301</b>	<b>Moved Permanently</b>	<b>The URI of the resource has changed</b>
<b>302</b>	<b>Found</b>	<b>The URI of the resource has temporarily changed</b>
303	<i>See Other</i>	Redirects client to a new resource requested at another address
304	<i>Not Modified</i>	Response has not been modified (caching purposes)
305	<i>Use Proxy</i>	Must be accessed from a proxy (deprecated)
306	<i>Temporary Redirect</i>	Redirects the client to another URI with the same method as the origin request

# HTTP/1.1 >> Response message (II)

**Status-Line** = HTTP-Version **Status-Code Reason-Phrase** CRLF

Code	Reason-phrase	Description
<b>CLIENT ERROR - 4XX (400-417)</b>		
<b>400</b>	<b><i>Bad request</i></b>	<b>Invalid request syntax</b>
<b>401</b>	<b><i>Unauthorized</i></b>	<b>Requires authentication to get the requested response</b>
402	<i>Payment Required</i>	Reserved for future use
<b>403</b>	<b><i>Forbidden</i></b>	<b>Customer does not have the required permissions</b>
<b>404</b>	<b><i>Not Found</i></b>	<b>The server could not find the requested resource</b>
405	<i>Method Not Allowed</i>	The method cannot be used
407	<i>Proxy Authentication Required</i>	Authentication must be done from a proxy (401)
408	<i>Request Time-out</i>	Inactive connection
413	<i>Request Entity Too Large</i>	The request entity is longer than the defined limits.
414	<i>Request-URI Too Large</i>	The URI is too large
415	<i>Unsupported Media Type</i>	Server does not support the media format of the requested data
<b>SERVER ERROR – 5XX</b>		
<b>500</b>	<b><i>Internal Server Error</i></b>	<b>Internal server error</b>
<b>501</b>	<b><i>Not Implemented</i></b>	<b>The requested method is not supported by the server</b>
<b>502</b>	<b><i>Bad Gateway</i></b>	<b>The server, as gateway, got an invalid response</b>
503	<i>Service Unavailable</i>	The server is not ready to handle the request
504	<i>Gateway Time-out</i>	The server, as a gateway, cannot have a response in time
<b>505</b>	<b><i>HTTP Version not supported</i></b>	<b>The HTTP version used in the request is not supported.</b>

## HTTP/1.1 >> Message headers

- Different types of headers are defined  $\Rightarrow$  header: value
  - **General headers**  $\Rightarrow$  apply to both requests and responses, unrelated to the data being transmitted in the body.
  - **Request headers**
  - **Response headers**
  - **Entity headers**  $\Rightarrow$  Define meta-information about the content.

### Some general headers

Cache-Control	Cache-Control: no-cache   Cache-Control: max-age=0 Cache-Control: public Cache-Control: private, community="UC3M"
Connection	Connection: close   Connection: upgrade
Date	Date: Sun, 04 Nov 2022 08:49:37 GMT
Pragma	Pragma: no-cache
Transfer-Encoding	Transfer-Encoding: chunked Transfer-Encoding: gzip
Upgrade	Upgrade: HTTP/2.0, HTTP/3.0   Upgrade: websockets
Via	Via: 1.0 fred, 1.1 nowhere.com (Apache/1.1)

"|" - in the examples - indicates a separation of different and independent headers with the possible values that can be defined

## HTTP/1.1 >> General headers >> Cache-Control

- `Cache-Control = "Cache-Control" ":" 1#cache-directive`
- `cache-directive = cache-request-directive | cache-response-directive`
- `cache-request-directive = "no-cache" [ "=" <"> 1#field-name <"> ]`
  - | `"no-store"`
  - | `"max-age" "=" delta-seconds`
  - | `"max-stale" [ "=" delta-seconds ]`
  - | `"min-fresh" "=" delta-seconds`
  - | `"only-if-cached"`
  - | `cache-extension`
- `cache-response-directive = "public"`
  - | `"private" [ "=" <"> 1#field-name <"> ]`
  - | `"no-cache" [ "=" <"> 1#field-name <"> ]`
  - | `"no-store"`
  - | `"no-transform"`
  - | `"must-revalidate"`
  - | `"proxy-revalidate"`
  - | `"max-age" "=" delta-seconds`
  - | `cache-extension`

## Cache-control options

Option	Description
<i>max-age=[seconds]</i>	Number of seconds that the resource is considered valid
<i>s-maxage=[seconds]</i>	The same as max-age but only for public/shared caches
<i>public</i>	Marks the resource as cacheable by any cache
<i>private</i>	Marks the resource as cacheable except for public/shared caches
<i>no-cache</i>	Revalidate the resource with the server before it serves it from its cache
<i>no-store</i>	The cache must not store a copy of the resource
<i>must-revalidate</i>	Cache must revalidate expired resource. i.e not serve if expired
<i>proxy-revalidate</i>	Same as must-revalidate header excluding private caches
<i>no-transform</i>	Instructs cache not to transform body of message

## HTTP/1.1 &gt;&gt; General headers &gt;&gt; Fragmented encoding

- Transfer-encoding: chunked means fragmented encoding, so the body of a message is sent as a series of *chunks*.
  - allows a client/server to send data without knowing in advance the total length of data to be sent (typical in dynamic pages)
- Chunks are sent:
  - Line with size in bytes (encoded in hex), followed by ";"
    - after the size there may be optional parameters that can be ignored
  - The chunk ends with <CRLF>
  - To indicate the end of the data, a line with "0" is sent
  - More headers (optional)
  - <CRLF>
- Optionally can include at the end of the general header Trailer

```
HTTP/1.1 200 OK
Date: Fri, 31 Dec 2012 23:59:59 GMT
Content-Type: text/plain
Transfer-Encoding: chunked
1a; ignore this parameter
abcdefghijklmnopqrstuvwxy
10
1234567890abcdef
0
some-header: some-value
another-header: another-value
[this is a blank line] (<CRLF>)
```



## HTTP/1.1 &gt;&gt; Request headers

Header	Example
<b>Accept</b>   <b>Accept-Charset</b>   <b>Accept-Encoding</b>   <b>Accept-Language</b>	<b>Accept:</b> text/plain; q=0.5, text/html, text/x-dvi; q=0.8, text/x-c <b>Accept:</b> application/json, text/json <b>Accept-Charset:</b> "iso-8859-5"
<b>Authorization</b>   Proxy-Authorization	<b>Authorization:</b> Basic A29X3DF13KLGF2U=
From	From: webmaster@w3.org
<b>Host</b> [mandatory]	<b>Host:</b> <u>www.w3.org</u> <b>Host:</b> <u>www.example.com:8080</u>
If-Match   If-None-Match   If-Range   If-Modified-Since   If-Unmodified-Since	Conditional requests If-Modified-Since: Sat, 29 Oct 2022 19:43:31 GMT
Max-Forwards (TRACE, OPTIONS)	Max-forwards: 1
Range	Range: bytes=0-499, -500
<b>Referer</b>	Referer: <a href="http://www.w3.org/hypertext/view.html">http://www.w3.org/hypertext/view.html</a>
TE	TE: deflate   TE: trailers
<b>User-Agent</b>	User-Agent: Mozilla/5.0

## HTTP/1.1 &gt;&gt; Request headers &gt;&gt; Authorization

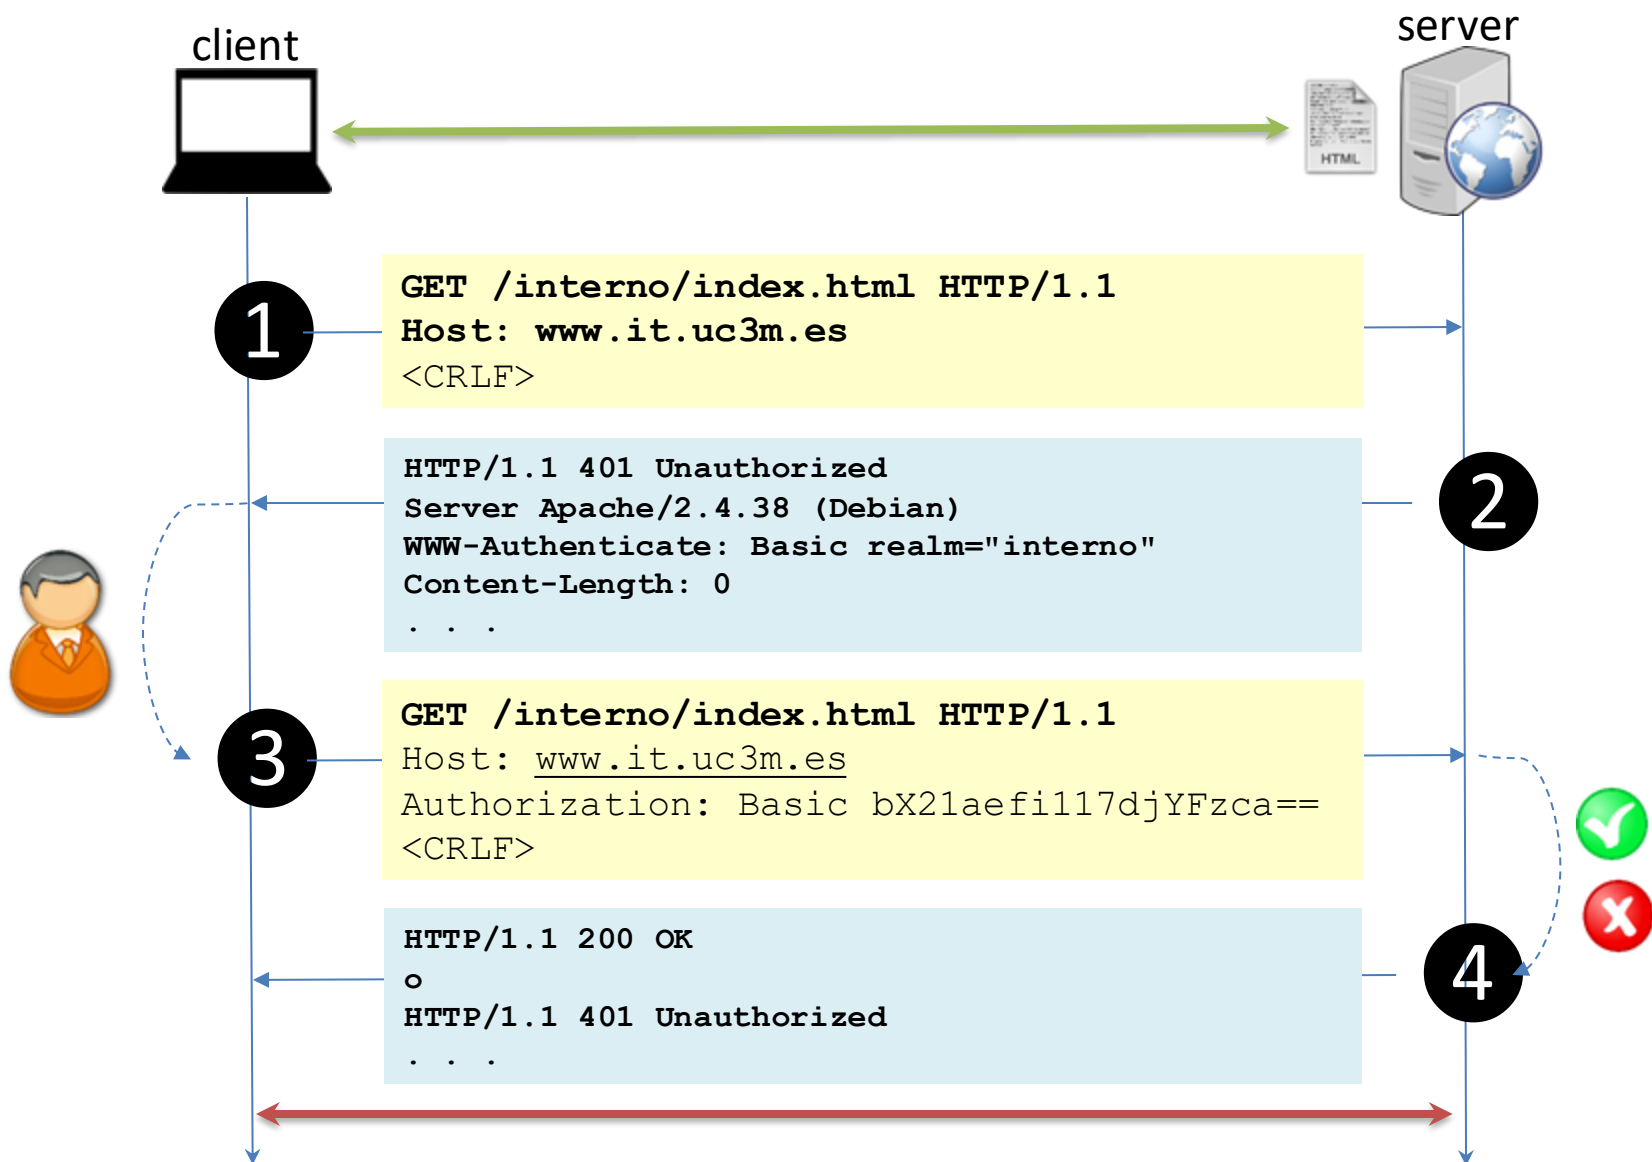
- Allows sending credentials to authenticate a user:
  - **Basic** authentication schema
  - **Digest** uses the value of a “nonce”
- The server can define different security zones (“realms”).
  - If a client requests a resource in one of this zones, a challenge is sent back in the response:

```
401 Unauthorized
WWW-Authenticate: Basic/Digest realm="interno"
[...]
```
- The client sends back the same request but includes a header with the security type, realm, username and password.

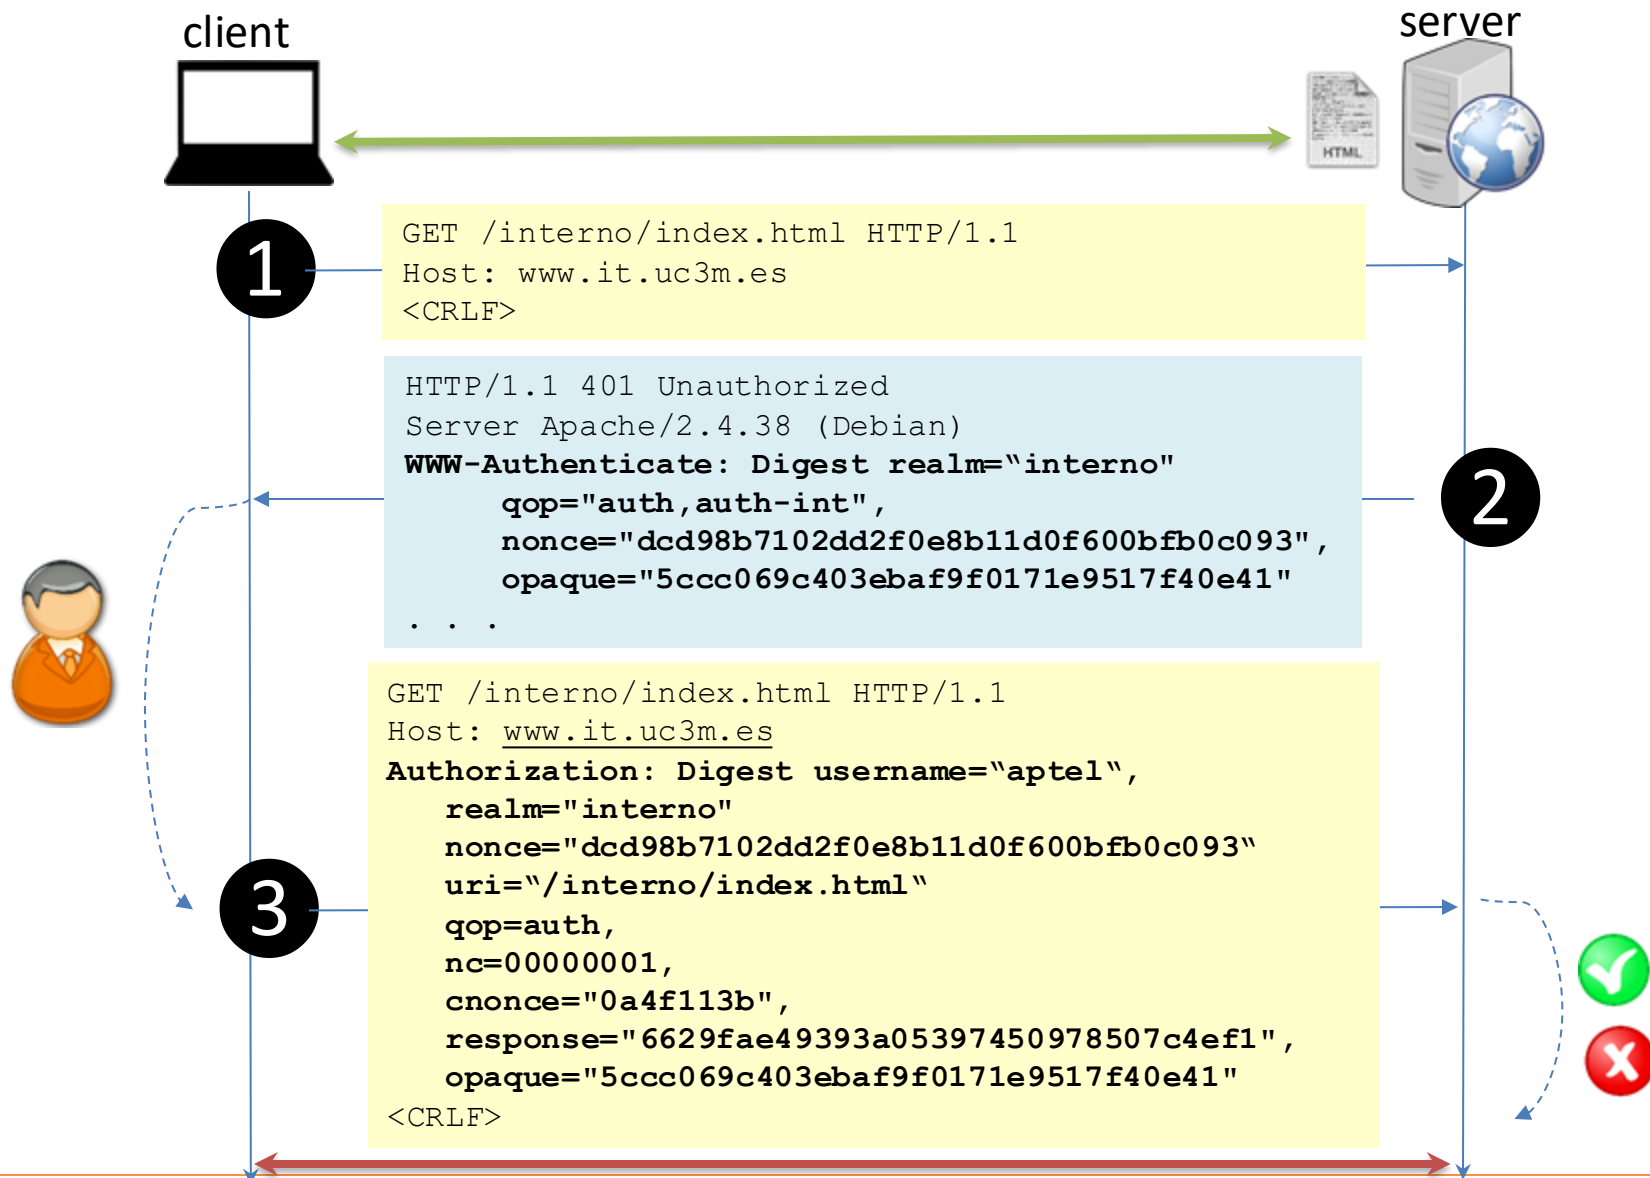
## HTTP/1.1 >> Types of authentication

- They are specified in RFC 2617 (updated in RFC 7616 and 7617)
- In the **basic** security scheme user:password is base64-encoded.
  - `Authorization: Basic djpvcGVuIHNlc2FtZQ==`
    - No confidentiality, nor message integrity
    - A trusted channel is assumed https (SSL, TLS)
- The **Digest** authentication scheme uses a nonce and other parameters
  - SHA-512/256 hash function is used, but allows MD5 for backward compatibility: `hash(a1:nonce:a2)`, where  
a1=`hash("username:realm:password")` and  
a2=`hash("reqMethod:reqURI")`
  - prevents replay attacks and mitigates nonce count attacks
  - defines a quality of protection parameter ("qop") for entity integrity (request body, not headers)

## HTTP/1.1 &gt;&gt; Example basic authentication



## HTTP/1.1 &gt;&gt; Example Digest authentication



## HTTP/1.1 &gt;&gt; Sending content-heavy requests

**Bandwidth** optimisation  $\Rightarrow$  avoid that the client sends in a request (POST or PUT) a body of many bytes, which the server will not be able to process or satisfy the requirements.

- **Example request:**

```
POST /docs HTTP/1.1
Host: www.example.com
Content-Type: application/pdf
Content-Length: 99000
Expect: 100-continue
```

**Response:**

HTTP/1.1 100 Continue

PDF file content is sent

HTTP/1.1 200 OK  
(or the corresponding response)

- If not accepted, response 417 Expectation Failed
- -If the client does not receive a response after a while, send the body.
- Compatibility may vary between implementations
  - Not all implementations can support or correctly handle this header.

## HTTP/1.1 &gt;&gt; Request headers &gt;&gt; Host

- **Mandatory** in **HTTP/1.1** requests
  - if it does not appear the server responds with a *400 Bad Request* message.
- The value corresponds to the name of the server in the URL (even the port if different from 80)
  - It is a clue for the server to know to which machine the client tried to connect
  - Allows the same machine (IP) to serve different web sites (virtual hosts)
- For example, if we request the resource `file.html` through the following URL <http://w.host1.com:8080/path/file.html>, the request would be:

```
GET /path/file.html HTTP/1.1
Host: w.host1.com:8080<CRLF>
<CRLF>
```

## HTTP/1.1 >> Example of a request GET

```
GET /index.html HTTP/1.1
```

⇒ *Line of request*

```
Date: Thu, 17 Nov 2022 17:00:00 GMT
```

```
Connection: close
```

⇒ *General headers*

```
Host: www.server.com
```

⇒ *Request headers*

```
From: admin@foo.com
```

```
Accept: text/html; text/plain
```

```
User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64)
```

```
AppleWebKit/537.36 (KHTML, like Gecko) Chrome/105.0.0.0
```

```
Safari/537.36
```

⇒ *Body of the request  
(empty) <CRLF>*



## HTTP/1.1 >> Example of a request POST

**POST /form HTTP/1.1**

Date: Thu, 17 Nov 2022 17:00:00 GMT

Connection: close

⇒ *General headers*

**Host: www.server.com**

⇒ *Request headers*

Content-Type: application/x-www-form-urlencoded

Content-Length: 19

nombre=pepe&edad=25

⇒ *Body of the request*

## HTTP/1.1 >> Response headers

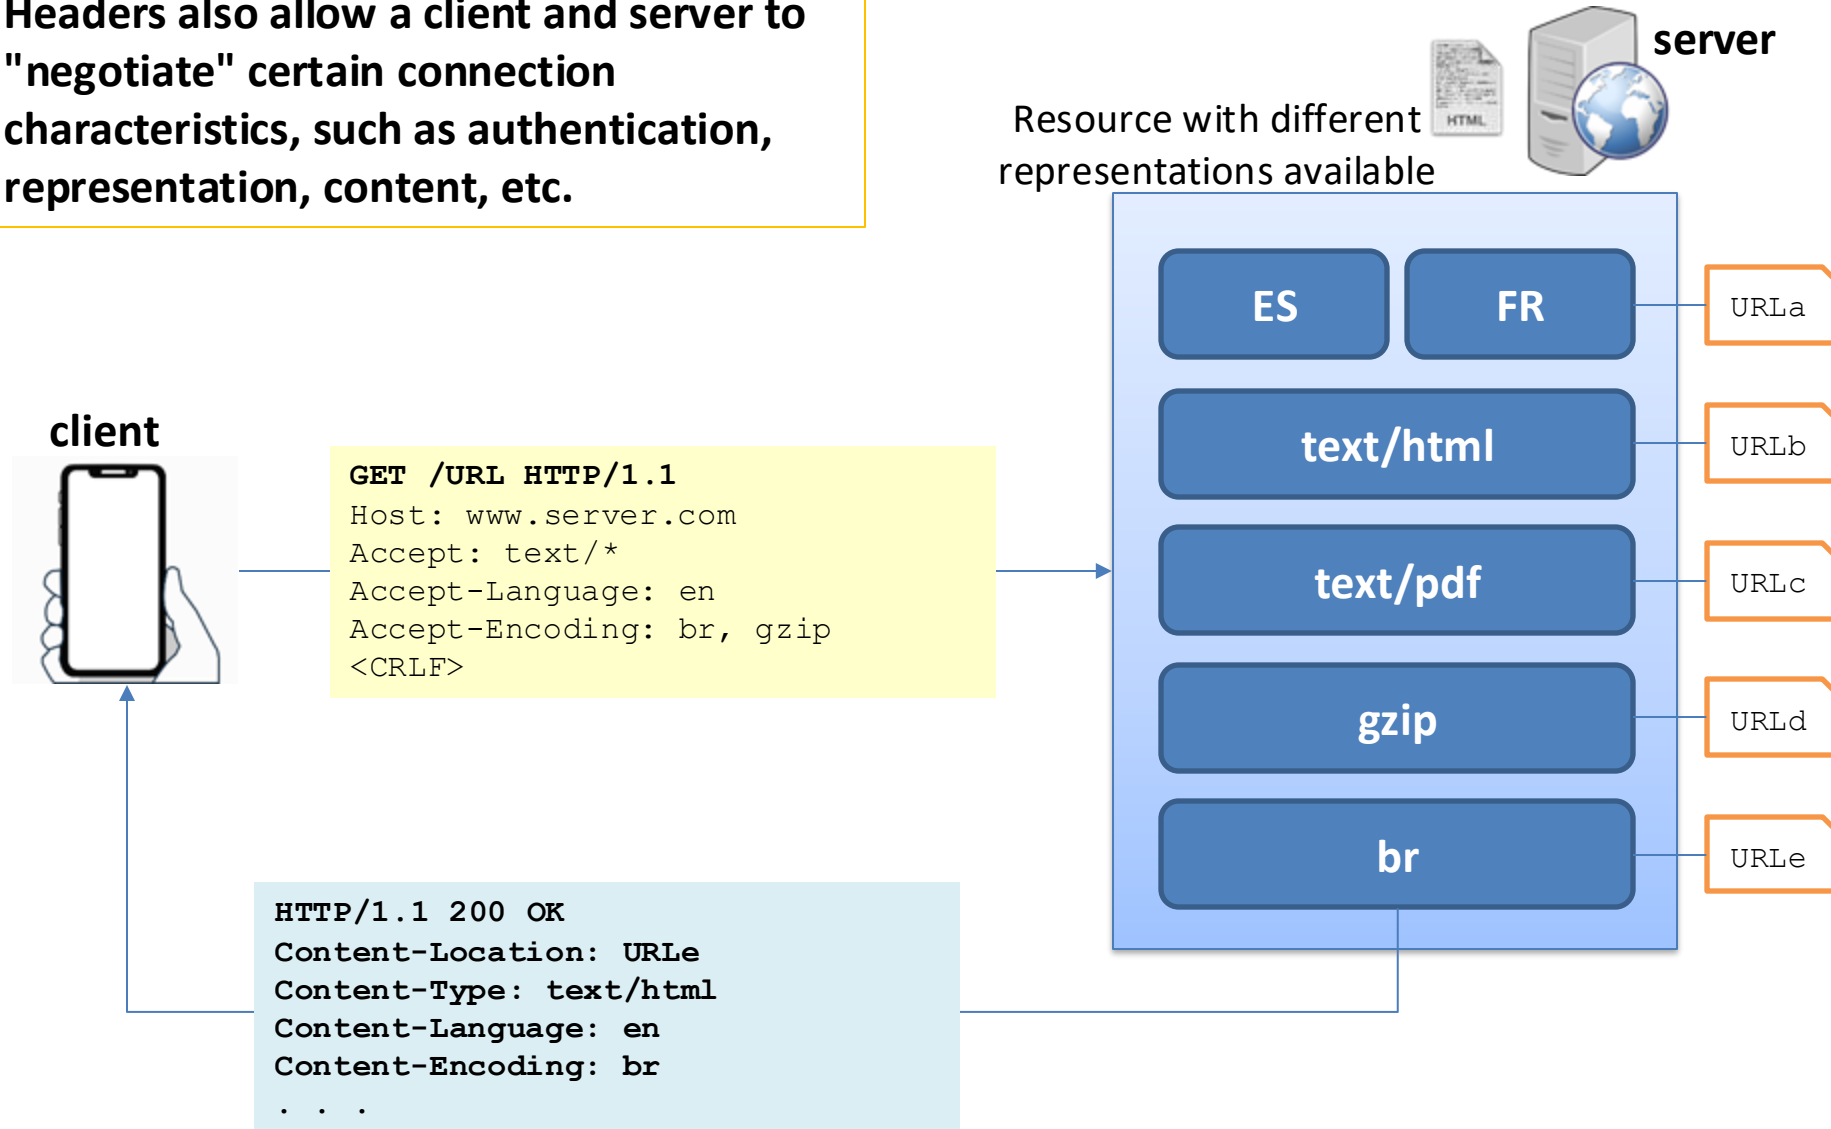
Header	Example
<b>Accept-Ranges</b>	Accept-Ranges: bytes   Accept-Ranges: none
Age	Age: 0   Age: 3600
<b>Etag</b>	ETag: W/"948-1501098468000"
<b>Location</b>	Location: <a href="http://www.w3.org/pub/people.html">http://www.w3.org/pub/people.html</a>
Proxy-Authenticate   <b>WWW-Authenticate</b>	WWW-Authenticate: Basic realm="interno"
Retry-After	Retry-After: Thu, 17 Nov 2022 23:00 GMT   Retry-After: 120
<b>Server</b>	Server: Apache/2.4.38 (Debian)
Vary	Vary: User-Agent   Vary: Accept-Encoding

## HTTP/1.1 >> Entity headers

Entity header	Example
Allow	Allow: GET, POST, HEAD
<b>Content-Encoding</b> Content-Language <b>Content-Length</b> Content-Location Content-MD5 Content-Range	Content-Encoding: gzip   Content-Encoding: br Content-Language: en-US Content-Length: 68137   Content-Length: 948 Content-Location: /documents/foo.json Content-Range: bytes 200-1000/67589
<b>Content-Type</b>	Content-Type: text/html; charset="ISO-8859-1" Content-Type: text/html; charset=utf-8 Content-Type: multipart/form-data; boundary=some
Expires	Expires: Sat, 31 Oct 2022 23:59:00 GMT
<b>Last-Modified</b>	Last-Modified: Wed, 28 Jul 2021 19:47:48 GMT

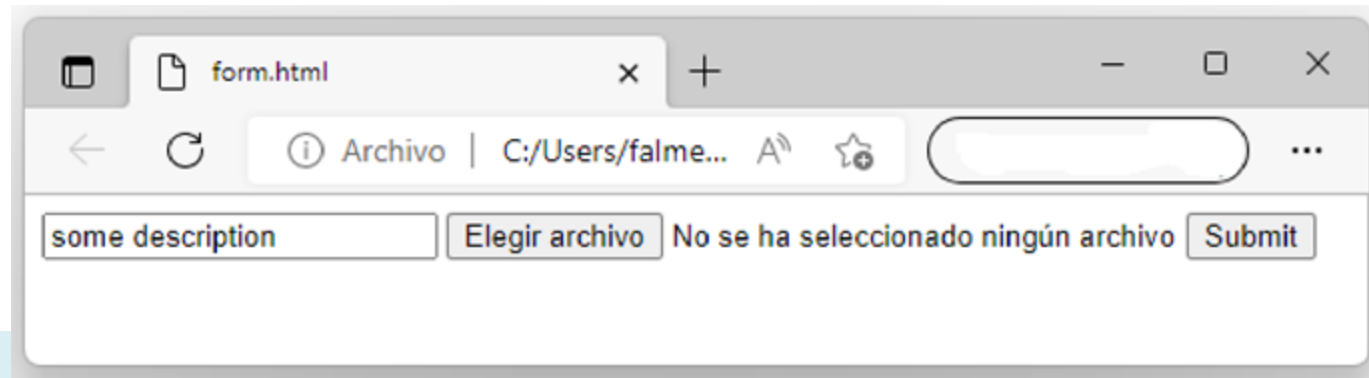
## HTTP/1.1 &gt;&gt; Entity headers &gt;&gt; Content negotiation

Headers also allow a client and server to "negotiate" certain connection characteristics, such as authentication, representation, content, etc.



## HTTP/1.1 &gt;&gt; Entity headers &gt;&gt; Content-Type: multipart

```
<form action="/" method="post" enctype="multipart/form-data">
  <input type="text" name="description" value="some description">
  <input type="file" name="someFile">
  <button type="submit">Submit</button>
</form>
```



```
POST /foo HTTP/1.1
```

```
Content-Length: 68137
```

```
Content-Type: multipart/form-data; boundary=-----859678012
```

```
-----859678012
```

```
Content-Disposition: form-data; name="description"
```

```
some description
```

```
-----859678012
```

```
Content-Disposition: form-data; name="someFile"; filename="example.txt"
```

```
Content-Type: text/plain
```

```
(content of the uploaded file example.txt)
```

```
-----859678012--
```

## HTTP/1.1 >> Response example

**HTTP/1.1 200 OK**

⇒ *Status line*

Date: Thu, 17 Nov 2022 17:01:00 GMT

Connection: close

⇒ *General headers*

Server: Apache/2.4

Accept-Ranges: bytes

Etag: "5ccca0c574cbebc38f6e95063286d500"

⇒ *Response headers*

Content-Type: text/html

Content-length: 144

Last-Modified: Mon, 14 Nov 2022 11:30:37 GMT

⇒ *Entity headers*

<html>

<head>

<title>Web: HTTP</title>

</head>

<body>

<h1>HTTP Lab</h1>

<p>Aplicaciones Telemáticas: Servidores Web (HTTP)</p>

</body>

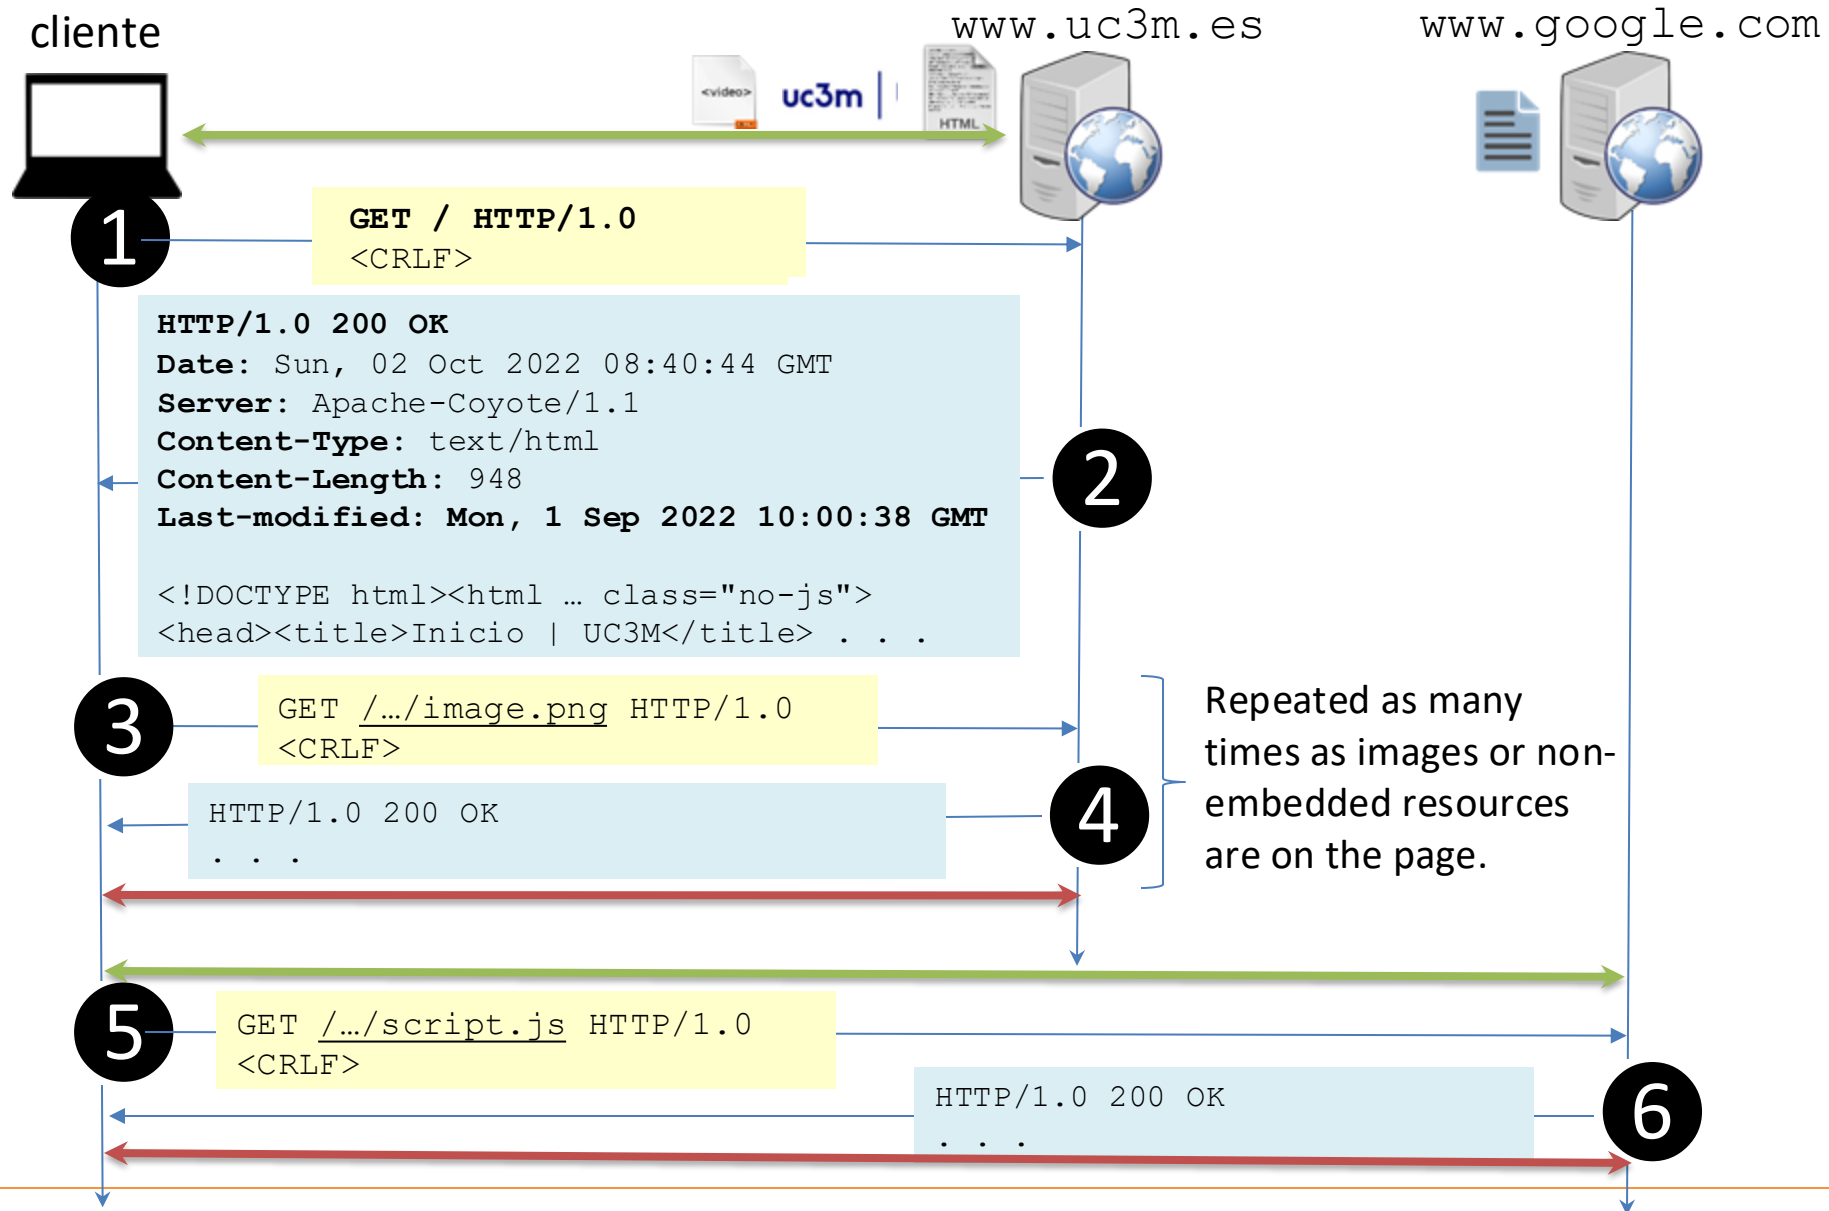
</html>

⇒ *Body of the response*

## HTTP/1.1 >> Persistent connections

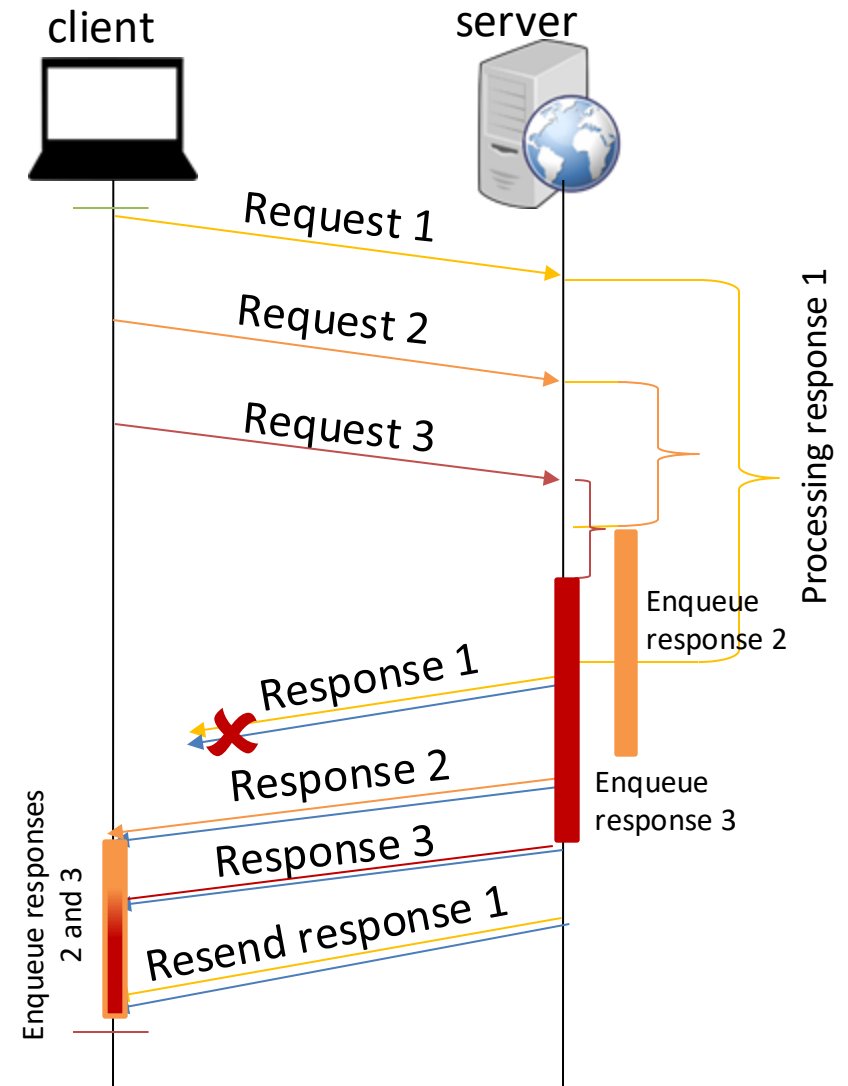

- Allows multiple requests to reuse the same connection
  - more than one transaction per request
  - several requests may be sent together without waiting for the responses (***HTTP pipelining***)
    - **responses** should be sent in the same order
    - the client must be prepared to resend the requests if the server close the connection before sending all the corresponding responses
  - The client can indicate that no more requests will be sent using the “`Connection: close`” header.
  - So the server may use the “`Connection: close`” header to close the connection (with or without sending all the responses)

## HTTP/1.1 >> Example persistent connections



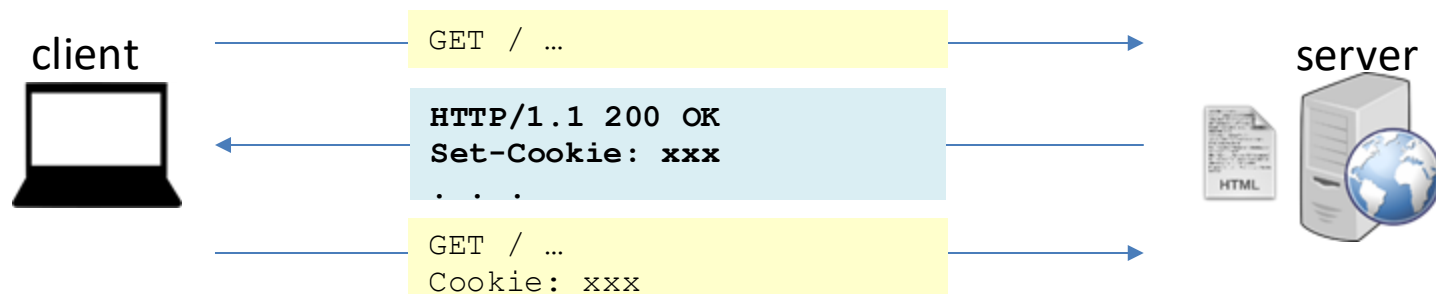


## HTTP/1.1 >> HOLB

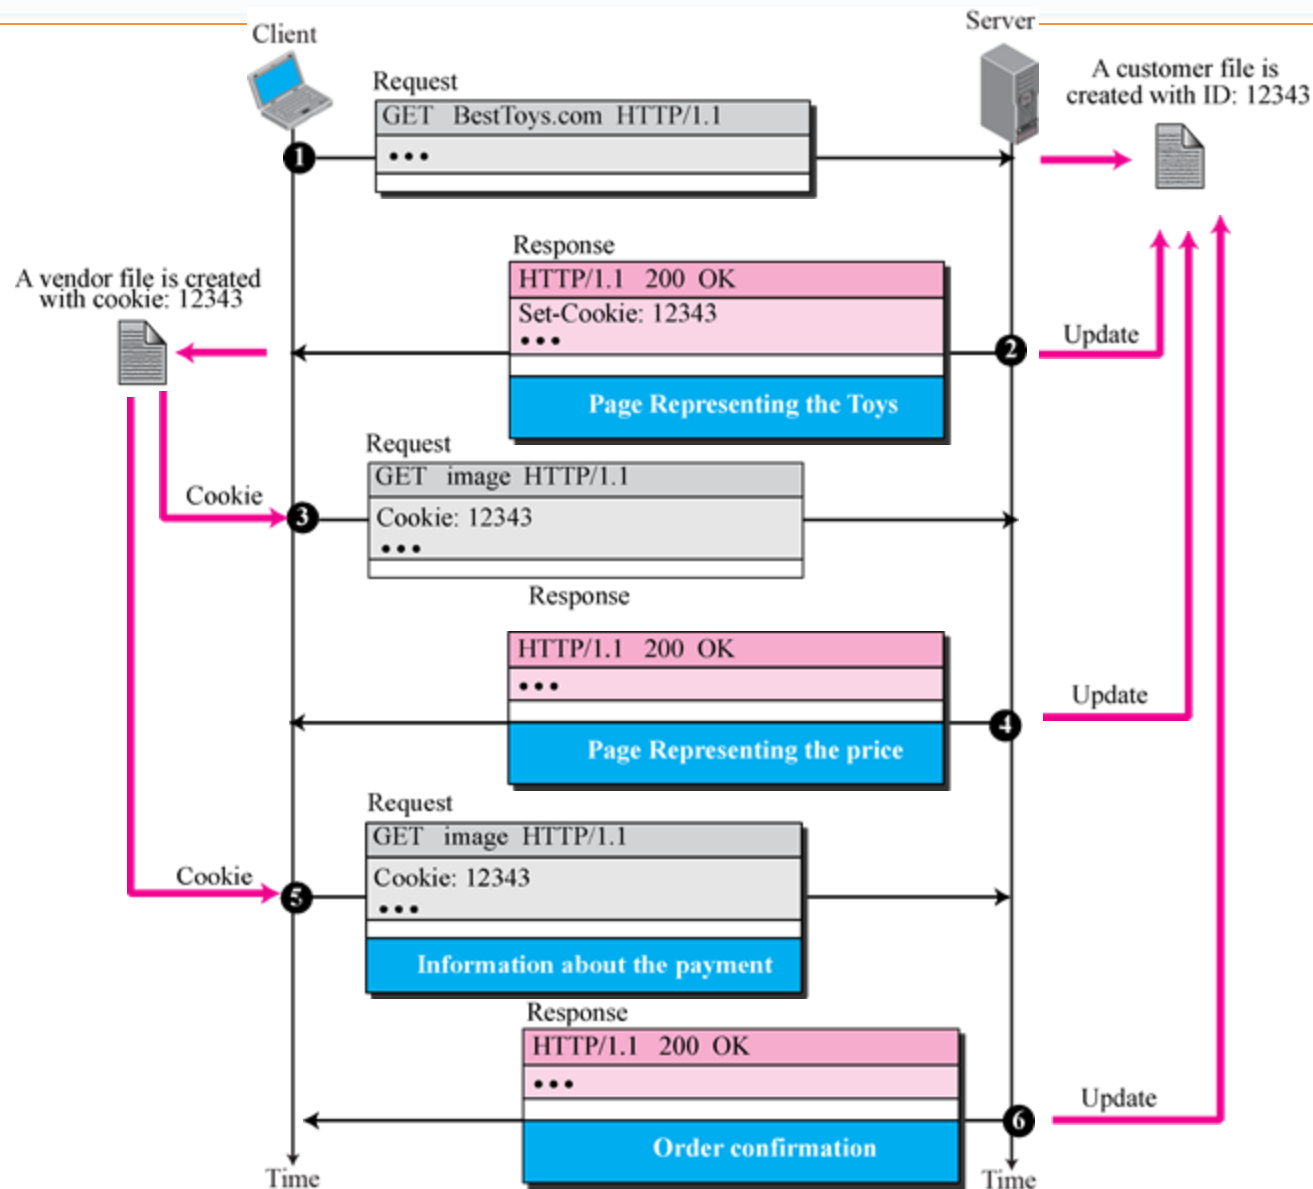


## HTTP/1.1 &gt;&gt; Status management &gt;&gt; Cookies

- Since HTTP is stateless, it can use **cookies**:
  - The server sends the cookie(s) with the HTTP response and the client stores them (header **Set-Cookie**).
  - Future requests from the client will be accompanied by the cookie(s) so that the server recognizes the client (and the session).
- Cookies were introduced by Netscape in 1994 (RFC 2965, updated with **RFC 6265**)
- A cookie allows to store context information about the session, user state, etc.
  - Browser does not interpret them, merely returns them to the server in future requests.
  - The client can impose restrictions or limits, e.g. number of cookies per server, totals, minimum length...



## HTTP/1.1 &gt;&gt; Status management &gt;&gt; Cookies



## HTTP/1.1. >> Exercises >> requests and responses

- Write the HTTP/1.1 request to retrieve an image with the URI `/www/image1.png` on the server [www.it.uc3m.es](http://www.it.uc3m.es)



- Which header do we add to indicate that the client can accept images in GIF, JPEG and PNG format?



- Write the answer defining the headers `Date`, `Server`, `Content-Type`, `Content-Length` of the image. Note: you do not need to include the body of the response.



## Content >> HTTP/2

### 4. HTTP/2

- A. Introduction & Background
- B. General overview
- C. Frame and frame types
- D. Requests and responses
- E. Exercise
- F. Compression and decompression
- G. Streams
- H. “PUSH”
- I. Conclusions

### Index

- 1. Introduction
- 2. First HTTP versions
- 3. HTTP 1.1
- 4. HTTP/2**
- 5. Other protocols

## HTTP/2 >> Differences HTTP/1.1 vs HTTP/2

From HTTP/2 IN ACTION by BARRY POLLARD, Copyright 2018.

© <https://freecontent.manning.com/animation-http-1-1-vs-http-2-vs-http-2-with-push/>

## HTTP/2 >> Differences HTTP/1.1 vs HTTP/2

From HTTP/2 IN ACTION by BARRY POLLARD, Copyright 2018.

© <https://freecontent.manning.com/animation-http-1-1-vs-http-2-vs-http-2-with-push/>



**HTTP/2 is the future of the Web, and it is here!**

**Your browser supports HTTP/2!**

This is a demo of HTTP/2's impact on your download of many small files making up the [Akamai Spinning Globe](#).





## HTTP/2 >> Introduction

- HTTP/2 was initially defined in **RFC 7540** (May 2015) and updated in **RFC 9113**, June 2022
  - derived from the open communication protocol that was developed for transporting web content, called **SPDY**, approx. 2012
- **SPDY**, protocol announced by Google in 2009 and developed in 2010
  - Purpose: reduce the loading time of web pages and improve web security
  - Main features:
    - Uses TCP and TLS
    - Allows **concurrent requests** in a single TCP session
      - client can assign a **priority** to each request, to prevent network congestion with non-critical resources
      - client keeps the **connection open** until it navigates away from the web pages or the server closes the connection

## HTTP/2 &gt;&gt; Background &gt;&gt; SPDY

- Main features of SPDY:
  - Reduces bandwidth (BW) by **compressing headers** (zlib) and eliminates unnecessary headers.
  - Defines message formats that are easy to parse ("**framing**")
    - “*control frames*” and “*data frames*”
  - Introduces **streams** that are independent bidirectional data streams divided into **frames**
    - 3 types of control frames to manage the stream lifecycle
  - Allows the server to initiate communications with the client and send data wherever possible
    - via a *stream* containing an “associated-stream-id”, which indicates the request *stream* to which the sent data is related.

## HTTP/2 >> General overview

- Optimize HTTP semantics, while maintaining backward compatibility (methods, URIs, codes, TCP/TLS...).
- Enable a more **efficient** use of network **resources**
  - clients establish a **single** network **connection** with the server
  - headers, data, messages are packaged in **binary data** structures
- Reduce **latency** by introducing header **compression** and allowing concurrent exchanges over the same connection
  - **multiplexing** and **prioritization** of requests  $\Rightarrow$  requests and responses associated with their own "**stream**".
  - avoid connection congestion and the problem of **head of line blocking** (HOLB) at HTTP level, i.e:
    - clients must limit the number of simultaneous connections
    - each client can make several requests on the same connection (HTTP/1.1) and the responses must be received in the same order

## HTTP/2 &gt;&gt; General overview

- Optimize HTTP semantics, while maintaining backward compatibility (methods, URIs, codes, TCP/TLS...).
- Enable a more **efficient** use of network **resources**
  - clients establish a **single** network **connection** with the server
  - headers, data, messages are packaged in **binary data** structures
- Reduce **latency** by introducing **header compression** and allowing concurrent exchange

- **multiplexing** responses and requests
- avoid connection **blocking** (HTTP/1.1)

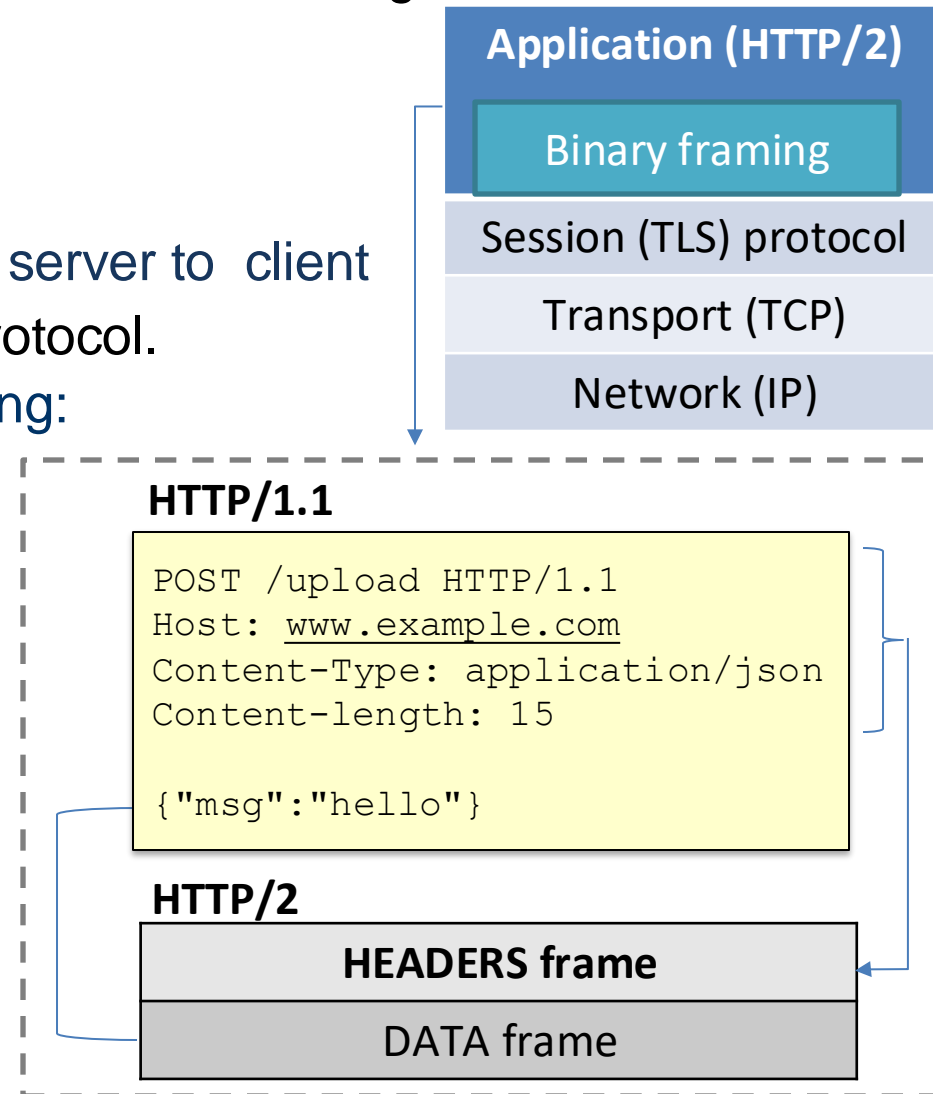
*“Clients that use persistent connections SHOULD limit the number of simultaneous connections that they maintain to a given server. A single-user client SHOULD NOT maintain more than 2 connections with any server or proxy... These guidelines are intended to improve HTTP response times and avoid congestion”.*  
[RFC2616]

- clients must limit the number of simultaneous connections
- each client can make several requests on the same connection (HTTP/1.1) and the responses must be received in the same order

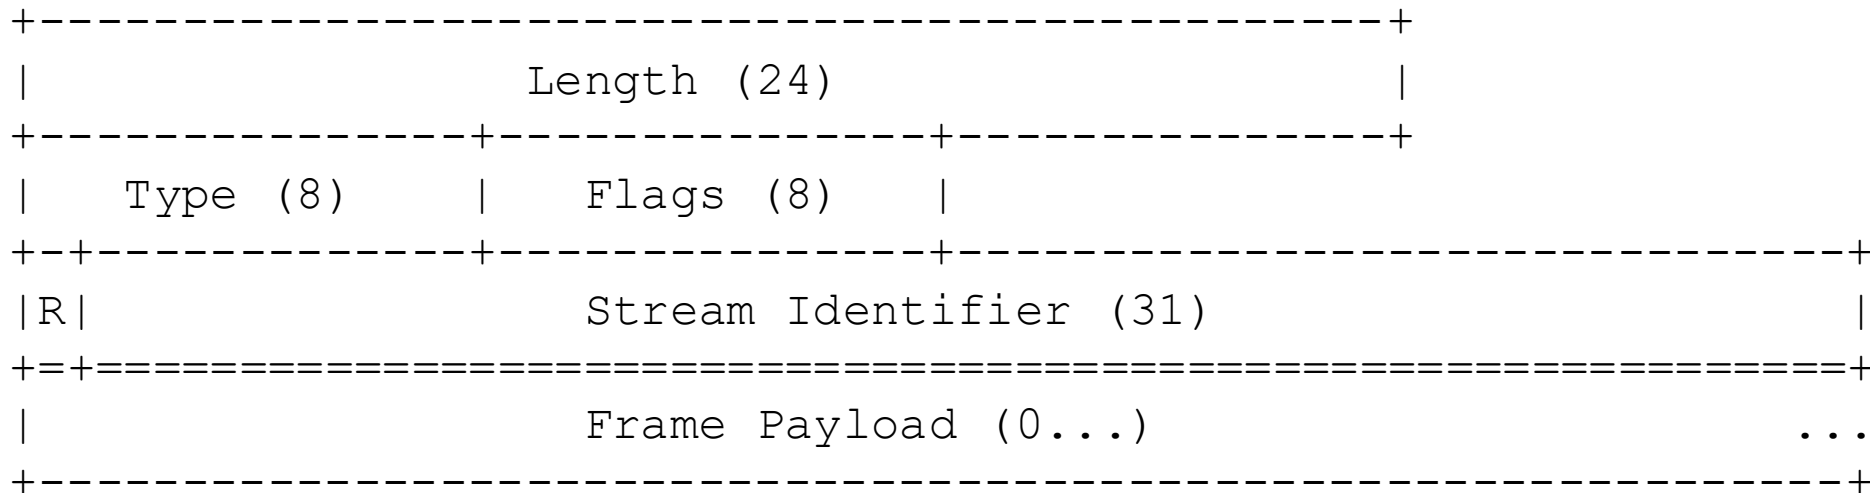
## HTTP/2 &gt;&gt; General overview (II)

- It can be generalized into two parts: the framing layer and the data or http layer.
- The binary framing layer
  - prioritization and flow control
  - introduce *push* messages from server to client
- **Frames** are the basic unit of the protocol.
  - have 9-octets headers including:

```
HTTP Frame {  
    Length (24),  
    Type (8),  
  
    Flags (8),  
    Reserved (1),  
    Stream Identifier (31),  
  
    Frame Payload (0...),  
}
```



## HTTP/2 &gt;&gt; Frame header fields



## Estructura de un "Frame"© RFC 7540

Field	Length	Description
<b>Length</b>	3 bytes	Payload length ( $2^{14}$ and $2^{24}$ -1 bytes)
<b>Type</b>	1 byte	<i>Frame</i> type
<b>Flags</b>	1 byte	Flags specific to the frame type
<b>R</b>	1 bit	Reserved
<b>Stream ID</b>	31 bits	Unique identifier for each stream
<b>Payload</b>	Variable	Data according to length specified in field Length

## HTTP/2 >> Frame types

Frame	Functionality
DATA	Carry octets in a stream, for instance HTTP requests or responses
HEADERS	To open a stream, or carry header block fragment
PRIORITY	Sender advised priority of a stream (carries dependency and weight)
RST_STREAM	Requires immediate termination of a stream
SETTINGS	Configuration parameters like preferences and constraints on peer behavior (HEADER_TABLE_SIZE, ENABLE_PUSH, MAX_CONCURRENT_STREAMS, INITIAL_WINDOW_SIZE, MAX_FRAME_SIZE, MAX_HEADER_LIST_SIZE)
PUSH_PROMISE	Notifies the peer endpoint (client or server) in advance of streams the sender intends to initiate
PING	Measuring round trip time
GOAWAY	To initiate shutdown of a connection or to signal serious error conditions
WINDOW_UPDATE	To implement flow control, signals of window increments up to $2^{31}-1$ , that may apply to a given stream or to the whole connection (stream_id=0)
CONTINUATION	To continue a sequence of header block fragments, the END_HEADERS flag is set if this is the last fragment

## HTTP/2 >> Connection management

- Connections are persistent
  - for performance, it is expected that clients will not close connections until it is determined that no further communication with a server is necessary (for example, when a user goes away from a particular web page) or until the server closes the connection.
  - Clients should not open more than one connection to a given “host:port” with the same configuration or under “normal” conditions.
- CONNECT method  $\Rightarrow$  converts an HTTP connection into a tunnel to a remote computer (used with proxies to establish a TLS session).
- HTTP/2 implementations must use TLS version 1.2 or higher when using a secure connection.
- Although the use of TLS is not mandatory in the specification, in practice most implementations (e.g., browsers) support HTTP/2 when used over TLS.



## HTTP/2 >> Request-Response

- A client making an HTTP request for a URI, without prior knowledge about HTTP/2 support, uses the `Upgrade` header
  - "h2c" for "http" connections
  - "h2" for "https" connections

```
GET / HTTP/1.1
Host: server.example.com
Connection: Upgrade, HTTP2-Settings
Upgrade: h2c
HTTP2-Settings: <base64url encoding
of HTTP/2 SETTINGS payload>
```

```
HTTP/1.1 200 OK
Content-Length: 243
Content-Type: text/html
...
```

```
HTTP/1.1 101 Switching Protocols
Connection: Upgrade
Upgrade: h2c

[ HTTP/2 connection ...
```

- If it has prior knowledge, then the client must send the preamble of the connection and immediately send frames (SETTINGS)

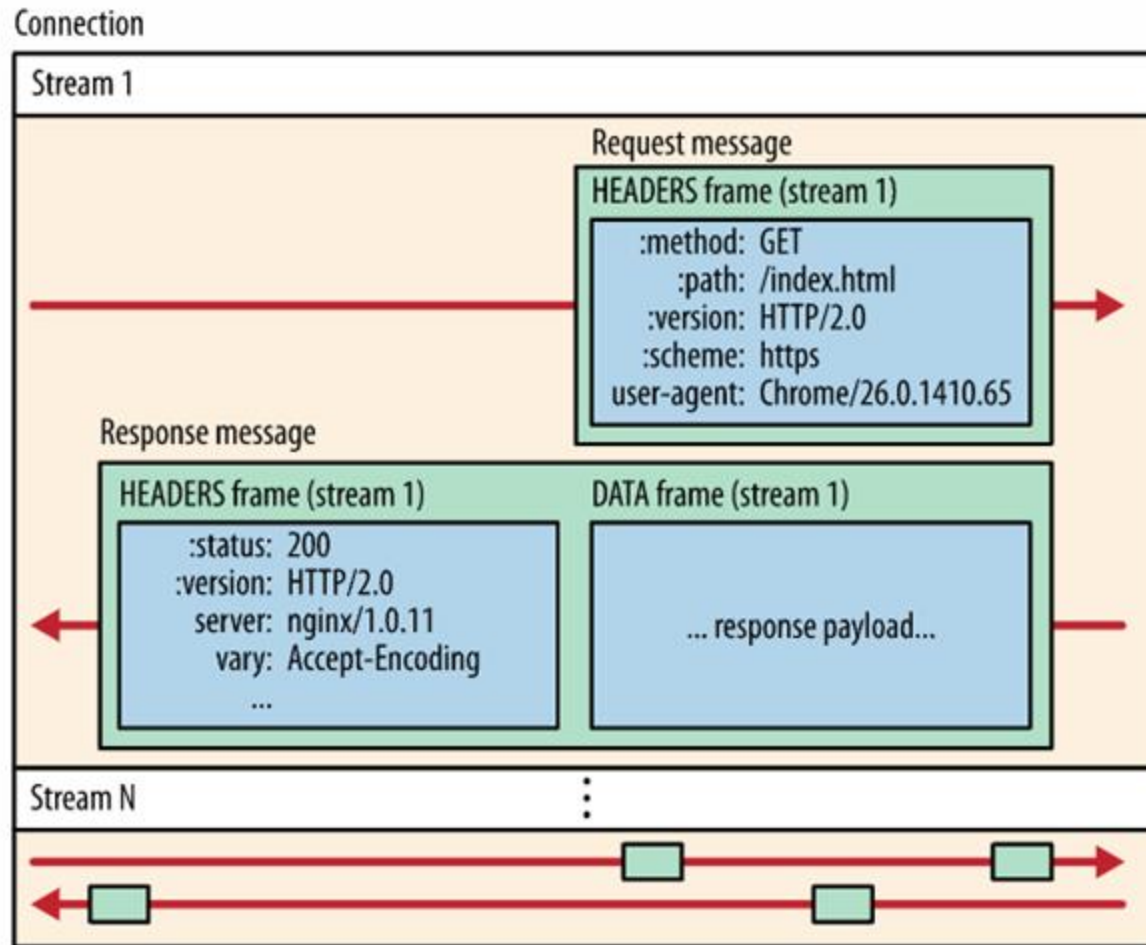
```
x505249202a20485454502f322e300d0a0d0a534d0d0a0d0a
```

HEX



```
"PRI * HTTP/2.0\r\n\r\nSM\r\n\r\n"
```

## HTTP/2 >>Request - response exchange



Source: <https://hpbn.co/http2/#streams-messages-and-frames>

## HTTP/2 >> Example of a GET request - response

```
GET /resource HTTP/1.1
Host: example.org
Accept: image/jpeg
```

==>

```
HEADERS
+ END_STREAM
+ END_HEADERS
:method = GET
:scheme = https
:path = /resource
host = example.org
accept = image/jpeg
```

```
HTTP/1.1 304 Not Modified
ETag: "xyzzzy"
Expires: Thu, 23 Jan ...
```

==>

```
HEADERS
+ END_STREAM
+ END_HEADERS
:status = 304
etag = "xyzzzy"
expires = Thu, 23 Jan ...
```

## HTTP/2 >> Example of a POST request - response

POST /resource HTTP/1.1

Host: example.org ==>

Content-Type: image/jpeg

Content-Length: 123

{binary data}

HEADERS

- END\_STREAM

- END\_HEADERS

:method = POST

:path = /resource

:scheme = https

CONTINUATION

+ END\_HEADERS

content-type = image/jpeg

host = example.org

content-length = 123

DATA

+ END\_STREAM

{binary data}

HTTP/1.1 200 OK

Content-Type: image/jpeg ==>

Content-Length: 123

{binary data}

length = 123

HEADERS

- END\_STREAM

+ END\_HEADERS

:status = 200

content-type = image/jpeg

content-

DATA

+ END STREAM

## HTTP/2 >> Practical example of an HTTP request-response

```
$ curl -v --http2 http://gitlab.gast.it.uc3m.es/aptel
* Trying 163.117.141.50:80...
* Connected to gitlab.gast.it.uc3m.es (163.117.141.50) port 80 (#0)
> GET /aptel HTTP/1.1
> Host: gitlab.gast.it.uc3m.es
> User-Agent: curl/7.88.1
> Accept: */*
> Connection: Upgrade, HTTP2-Settings
> Upgrade: h2c
> HTTP2-Settings: AAMAAABkAAQCAAAAAAIAAAAA
>
< HTTP/1.1 301 Moved Permanently
< Server: nginx
< Date: Wed, 15 Nov 2023 13:23:07 GMT
< Content-Type: text/html
< Content-Length: 178
< Connection: keep-alive
< Location: https://gitlab.gast.it.uc3m.es/aptel
<
<html><head><title>301 Moved Permanently</title></head>
<body bgcolor="white"><center><h1>301 Moved Permanently</h1></center>
<hr><center>nginx</center>
</body></html>
```

# HTTP/2 >> Practical example of an HTTPS request-response

```
$ curl -v --http2 https://gitlab.gast.it.uc3m.es/aptel
* Trying 163.117.141.50:443...
* Connected to gitlab.gast.it.uc3m.es (163.117.141.50) port 443 (#0)
* ALPN: offers h2,http/1.1
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* . . .
* SSL connection using TLSv1.3 / TLS_AES_256_GCM_SHA384
* ALPN: server accepted h2
* . . .
* Using Stream ID: 1 (easy handle 0x5630d704fc70)
> GET /aptel HTTP/2
> Host: gitlab.gast.it.uc3m.es
> user-agent: curl/7.88.1
> accept: */*
>
< HTTP/2 200
< server: nginx/1.14.2
< date: Wed, 15 Nov 2023 13:17:20 GMT
< content-type: text/html; charset=utf-8
< content-length: 32611
< cache-control: max-age=0, private, must-revalidate
< etag: W/"6e19af91479c8bbba703be9fa61cd46c"
< referrer-policy: strict-origin-when-cross-origin
< vary: Accept
<
<!DOCTYPE html>
<html class="with-header with-system-footer" lang="en">
<head prefix="og: http://ogp.me/ns#"> . . .
```

## HTTP/2 >> Exercise

- Access to `aulavirtual.lab.it.uc3m.es`
- Open a terminal
- Access to <https://www.google.com> using the command line `curl`. Note: you can change the “User-agent”, e.g, `--user-agent "Mozilla/5.0 (X11; Linux x86_64; rv:91.0) Gecko/20100101 Firefox/91.0"`
- Indicates the request and response obtained (without the body)



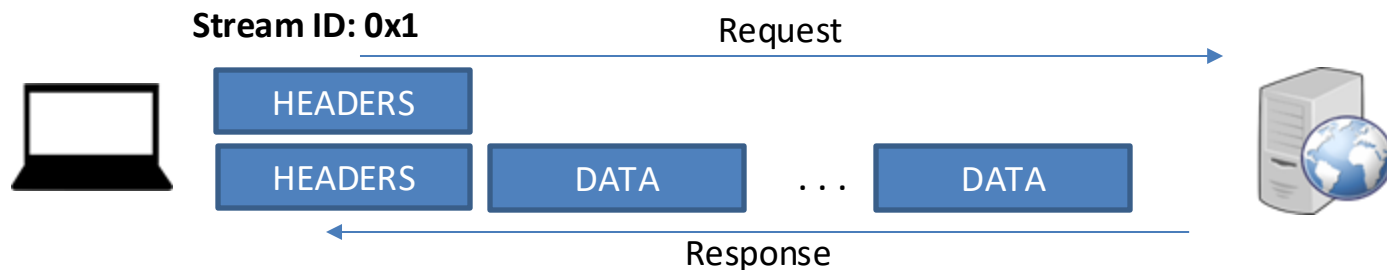
## HTTP/2 >> Compression

- HTTP message headers can contain large amounts of data, the “*frames*” contained in these messages are compressed (**HPACK**).
  - the main advantage lies in the size of the requests.
- Headers are used in both requests and responses in *push* messages.
- Header block fragments can be sent as *payload* of the frames HEADERS, PUSH\_PROMISE or CONTINUATION
  - who receives these frames organizes the header blocks and decompresses them.




## HTTP/2 &gt;&gt; Streams and Multiplexing

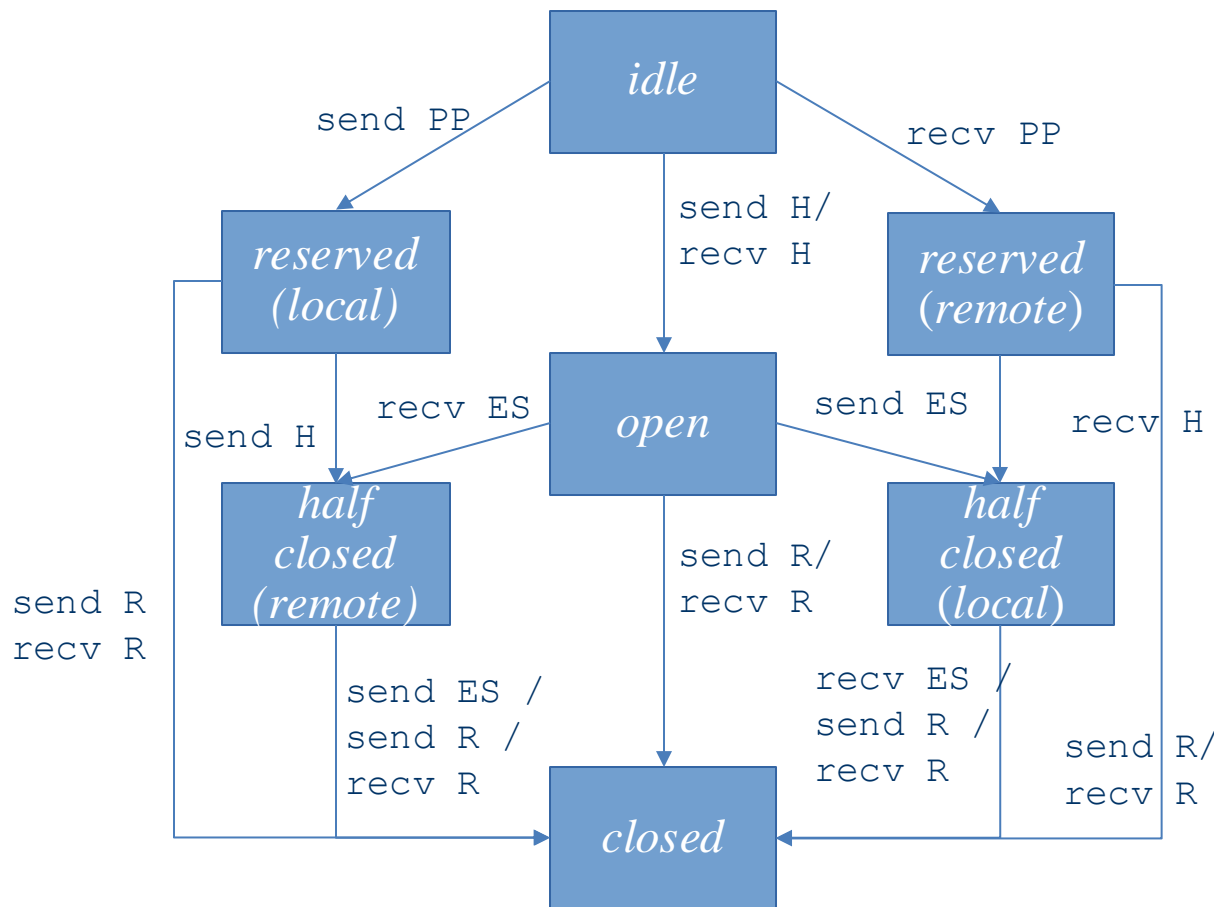
- A stream is a bidirectional, independent sequence of frames.



- One connection can have multiple streams open concurrently.
- Streams can be set up and used unilaterally or shared by client or server.
  - can be closed by either party
- The order in which the frames are sent is relevant:
  - processed in the order they are received
  - identified by an integer assigned by the initiator of the "stream"



# HTTP/2 >> Life cycle of a stream



H: HEADERS frame (with implied CONTINUATION)  
 PP: PUSH\_PROMISE frame (with implied CONTINUATION)  
 ES: END\_STREAM flag  
 R: RST\_STREAM frame

In the reserved state, WINDOW\_UPDATE or PRIORITY can be received, and PRIORITY can be sent

In the open state, any frame may be sent and received

In the half closed (local) state, frames of any type may be received

In the half closed (remote) state, frames of any type may be sent

## HTTP/2 >> Priority of "streams"

- A client can assign a priority for a new stream by including prioritization information in the HEADERS frames that open the stream.
- Later the stream priority can be changed with a PRIORITY frame
- **Priority**  $\Rightarrow$  How would you prefer the server (or the client) to allocate resources to this flow when managing concurrent streams? In other words, when the capacity to send is limited, which streams can transmit frames?
- Streams can be prioritized by marking them as dependent on the completion of other streams
  - a tree is formed with parent and child streams
- Each dependency is assigned a relative weight
- The processing of concurrent streams in a particular order cannot be forced using this priority. This is only a suggestion.

## HTTP/2 >> Error handling and connection closure

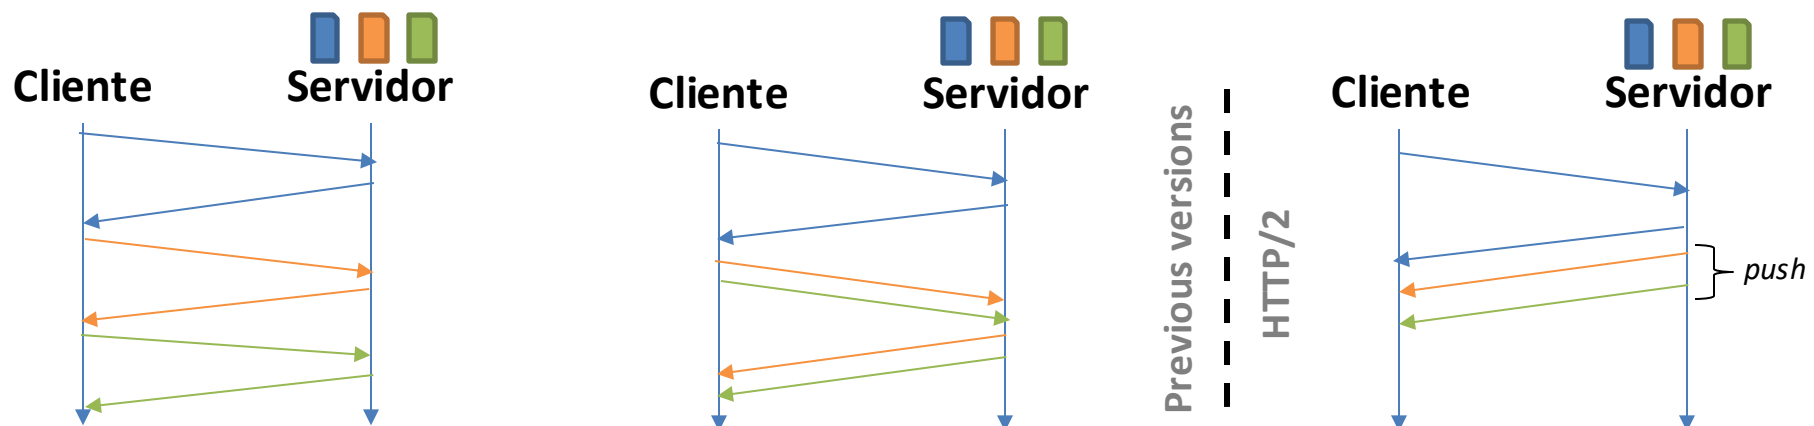
- Two types of errors can be handled:
  - Connection error  $\Rightarrow$  disables the entire connection
    - should send a GOAWAY frame with the “stream-id” of the last *stream* from which messages were successfully received.
    - After the transmission, the TCP connection must be closed.
  - Stream error  $\Rightarrow$  in an individual stream
    - Sends a RST\_STREAM frame with the “stream-id”, including an error code
- A list of error codes is defined indicating the cause, e.g., PROTOCOL\_ERROR, INTERNAL\_ERROR, SETTINGS\_TIMEOUT, REFUSED\_STREAM, HTTP\_1\_1\_REQUIRED...
- A connection can be closed by either party at any time, influencing the streams that are in state “*open*” or “*half-closed*”

## HTTP/2 >> PUSH from the server

- It is a new interaction model, where the **server** sends responses/data to a client without an explicit request ("**push**"), anticipating the client's needs
  - e.g., when the main resource requested includes other resources that are not embedded
  - therefore, the push is associated with a previous request initiated by the client.
- It uses a PUSH\_PROMISE frame.
- A client can disable that push (SETTINGS\_ENABLE\_PUSH)
  - it is negotiated at each hop independently.

## HTTP/2 &gt;&gt; PUSH from the server

- It is a new interaction model, where the **server** sends responses/data to a client without an explicit request ("**push**"), anticipating the client's needs
  - e.g., when the main resource requested includes other resources that are not embedded
  - therefore, the push is associated with a previous request initiated by the client.
- It uses a PUSH\_PROMISE frame.
- A client can disable that push (SETTINGS\_ENABLE\_PUSH)
  - it is negotiated at each hop independently.



## HTTP/2 >> Conclusions

- The most widely used browsers and servers have full HTTP/2 support for a few years now.
- HTTP/2 improves performance over HTTP/1.1, in terms of latency and flow control of TCP connections.
  - avoids HOLB at HTTP level
- However, HOLB at TCP level is still present.
  - **QUIC** protocol ("*Quick UDP Internet Connections*")
- **HTTP/3** has been defined in RFC 9114, September 2022, based on QUIC, instead of TCP/TLS.

## Content >> HTTP/2

### 4. Other protocols

A. QUIC

B. HTTP/3

### Index


1. Introduction
2. First HTTP versions
3. HTTP 1.1
4. HTTP/2
- 5. Other protocols**



## QUIC >> General overview

- Quick UDP Internet Connections emerged as an experimental protocol over the transport layer designed by J. Roskind (Google), 2013
  - submitted to the IETF for consideration in 2015
- Standardized by the IETF in May 2021, **RFC 9000**, as a protocol for **multiplexed and secure** connections over **UDP** (User Datagram Protocol).
- Similar to HTTP/2, but emerges as an alternative to TCP
- Designed to provide security equivalent to TLS, with lower connection and transport latency.
- It is enabled by default in Chromium and introduced in Opera 16.
- When the client accesses a web page, certificates and keys are sent.
- Client and server exchange QUIC packets, which contain **frames** with **control** information and application **data**.

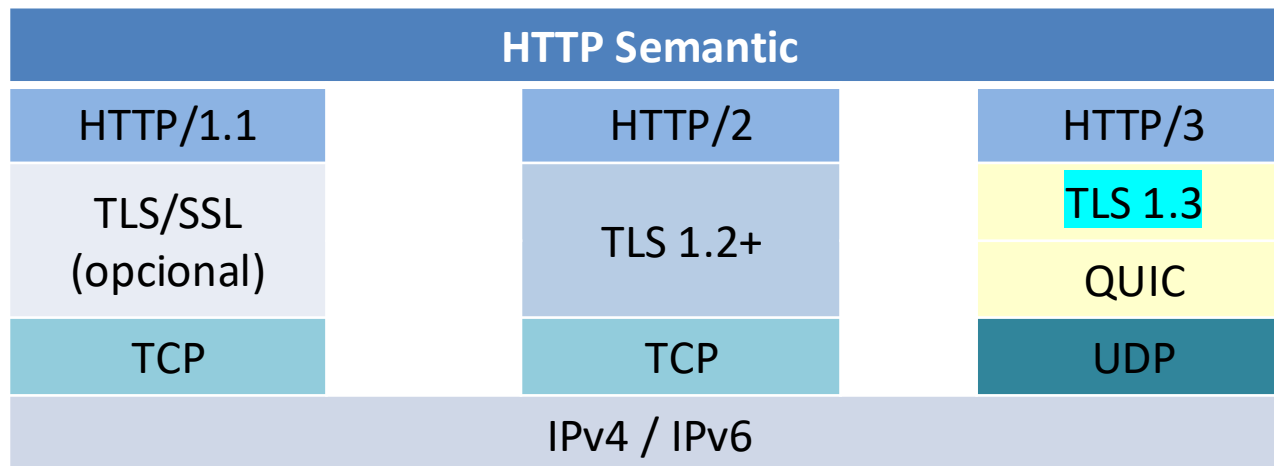
## QUIC >> General overview



Source: cloudflare.com

## HTTP/3 &gt;&gt; General overview

- **RFC 9114** describes how to map HTTP over QUIC, identifying features of HTTP/2 that are assumed by QUIC and how extensions can be transferred to HTTP/3.
- When a client knows that the server implements HTTP/3, it opens a QUIC connection.
  - QUIC provides protocol negotiation, stream-based multiplexing and flow control.



- A client can attempt to access a resource with "https", using the "h3" token in the Application-Layer Protocol Negotiation (ALPN, *Application-Layer Protocol Negotiation*) in the TLS handshake.

## HTTP/3 >> General overview (II)

- **Similarities** with **HTTP/2**: connection management, HTTP semantics
  - persistent connections across multiple requests
  - HTTP messages consist of: header section, content (optional) and "trailer" section
    - request fields: `:method`, `:scheme`, `:path`
    - response fields: `:status`
- A client sends an HTTP request on a request stream, which is a bidirectional QUIC stream initiated by the client.
  - a single request must be sent on a given stream
  - the server sends zero or more HTTP responses on the same stream, followed by a single final HTTP response
- Push messages are sent on a one-way QUIC stream initiated by the server.
- **Differences** with **HTTP/2**: does not support the "Upgrade" mechanism, does not provide a mechanism to indicate priority, compression type, frame format: `Type`, `Length`, `Payload` (..)

The screenshot shows a Google Chrome browser window with the Google homepage. The DevTools Network tab is open, displaying a list of network requests. The selected request is a GET request to the Google search endpoint. The request headers are visible, showing the authority, method, path, scheme, and various headers like accept, accept-encoding, and accept-language. The cookies are also visible, including the CONSENT cookie.